

MATH 285J: Topics in High-Dimensional Sampling and Generative Diffusions

Lecture Notes (please use carefully; any typos let us know!)

Yifan Chen

UCLA, Department of Mathematics

April 2026

Contents

1	Course Goals	4
2	Static Methods: Rejection and Importance Sampling	5
2.1	Generating Samples in One Dimension	5
2.2	Rejection Sampling	5
2.3	Importance Sampling	6
2.4	Asymptotic Theory	6
2.5	Non-Asymptotic Bound and the Role of χ^2	7
2.6	The Curse of Dimensionality	7
3	Probability Divergences	9
3.1	f -Divergences	9
3.2	Special Cases	9
3.3	Comparison and High-Dimensional Behavior	9
3.4	Preview: Density Evolution under ODE/SDE Dynamics	10
	Week 1 Exercises	11
4	Langevin Dynamics	13
4.1	The Overdamped Langevin SDE	13
4.2	Fokker–Planck Equation and Stationarity	13
4.3	Mixing Time and the Wasserstein Metric	14
4.4	Convergence under Strong Convexity	14
4.5	Initialization and Dimension Dependence	15

5	Langevin Monte Carlo	17
5.1	From Continuous Time to Algorithms	17
5.2	Warm-up: Gradient Descent for Optimization	17
5.3	ULA: Main Theorem	17
5.4	Proof of Theorem 5.1	18
5.5	Algorithmic Complexity	19
	Week 2 Exercises	21
6	Score-Based Diffusion Models	23
6.1	Motivation and Setting	23
6.2	Score Matching	23
6.3	Multiscale Approach and the Noising Process	24
6.4	Time Reversal and the Reverse SDE	25
6.5	Denoising Score Matching	25
6.6	Error Analysis	27
	Week 3 Exercises	30
7	Flow Matching, Rectified Flow, Stochastic Interpolants	32
7.1	Motivation: Beyond Gaussian Noise Source and Infinite Time	32
7.2	Setup and Interpolation Process	32
7.3	The Matching Theorem	33
7.4	Algorithm: Training by Matching the Velocity	33
7.5	Connection to Score Based Diffusion Models	34
7.6	Rectified Flow and Optimal Transport	36
	Week 4 Exercises	37
8	Conditional Sampling with Generative Priors	39
8.1	Motivation and Setting	39
8.2	Setting 1: Amortized Conditional Flow Matching	39
8.3	Setting 2: Guidance without Re-training	41
8.4	Discussions and Summary	42
	Week 5 Exercises	44
9	Stochastic Optimal Control for Conditional Sampling	47
9.1	Motivation	47
9.2	The Stochastic Optimal Control Problem	47
9.3	Value Function and Dynamic Programming	48
9.4	The Hamilton–Jacobi–Bellman Equation	48
9.5	Characterizing the Distributions via Girsanov’s Theorem	49
9.6	Connection to Sampling: KL-Control Duality	50
9.7	Pathwise Gradient and the Adjoint Equation	51

Week 6 Exercises	58
10 Sampling as Optimization in Probability Space	63
10.1 Langevin as KL Minimization	63

1 Course Goals

The central problem of this course is: *how do we sample from a target distribution $\pi \propto e^{-V}$ on $\mathbb{R}^d$ and compute statistics $\mathbb{E}_\pi[h]$, efficiently in high dimensions?* We assume access to π either via (1) **query access** to V or ∇V (with unknown normalization $Z = \int e^{-V} dx$), or (2) **data access** via existing samples $X_1, \dots, X_N \sim \pi$. Motivating settings include:

- **Bayesian inverse problems:** given $y = F(\theta) + \xi$ where $\xi \sim \mathcal{N}(0, \sigma^2 I_d)$ and a prior $p_{\text{prior}}(\theta)$, sample the posterior $p_{\text{post}}(\theta) \propto p_{\text{prior}}(\theta) e^{-\|y - F(\theta)\|^2 / (2\sigma^2)}$.
- **Statistical physics:** Gibbs measure $\pi \propto e^{-V/T}$ with potential V and temperature T ; partition function $Z = \int e^{-V/T} dx$ is intractable.
- **Generative modeling:** π is an unknown data distribution accessible only through samples.
- **Reinforcement learning:** V relates to reward functions.

Course outline and references

- **Static methods:** rejection/importance sampling, curse of dimensionality.
- **Dynamical methods:** MCMC, Langevin SDEs, gradient flow interpretation, mixing time and bias, acceleration by kinetic Langevin, preconditioning by affine invariant methods, de-bias via Metropolis-Hasting.
- **Flow and diffusion methods:** diffusion models, flow matching (ODEs/SDEs), posterior sampling with generative priors, stochastic optimal control, variational inference, and discrete diffusion.

References: Sanz-Alonso & Al-Ghattas (Introduction to Monte Carlo) [18]; Chewi (Algorithmic Complexity of Log-concave Sampling) [7]; Liu (Monte Carlo Strategies in Scientific Computing) [15]; Sanz-Alonso, Stuart & Taeb (Inverse Problems and Data Assimilation) [19]; various diffusion/flow papers.

2 Static Methods: Rejection and Importance Sampling

We start with the setting where we have query access V or its gradients. Namely, we know the density up to normalization constants.

2.1 Generating Samples in One Dimension

We assume throughout that $U \sim \text{Unif}(0, 1)$ can be generated. This suffices to sample any one-dimensional distribution via the following standard fact.

Proposition 2.1 (Inverse CDF method). *Let π be a probability distribution on $\mathbb{R}$ with CDF $F(x) = \mathbb{P}_\pi(X \leq x)$. If $U \sim \text{Unif}(0, 1)$, then $X = F^{-1}(U) \sim \pi$, where $F^{-1}(u) = \inf\{x : F(x) \geq u\}$.*

Proof. $\mathbb{P}(F^{-1}(U) \leq x) = \mathbb{P}(U \leq F(x)) = F(x)$. $\square$

This gives Gaussian, exponential, and other standard distributions. For general high-dimensional targets, the CDF is intractable and more powerful methods are needed.

2.2 Rejection Sampling

Setup. Let $\pi, g : \mathbb{R}^d \rightarrow [0, \infty)$ be probability densities with $\text{supp}(\pi) \subseteq \text{supp}(g)$, and fix $M \geq \sup_{x \in \mathbb{R}^d} \pi(x)/g(x) < \infty$.

Algorithm. Repeat until acceptance:

- (1) Draw $Y \sim g$ and $U \sim \text{Unif}(0, 1)$ independently.
- (2) Compute $r(Y) = \pi(Y)/(Mg(Y)) \in [0, 1]$.
- (3) If $U \leq r(Y)$, return $X = Y$. Otherwise go to (1).

Proposition 2.2 (Correctness). *The returned sample satisfies $X \sim \pi$.*

Proof. For any measurable $A \subseteq \mathbb{R}^d$,

$$\mathbb{P}(Y \in A \mid Y \text{ accepted}) = \frac{\int_A g(y) r(y) dy}{\int_{\mathbb{R}^d} g(y) r(y) dy} = \frac{\frac{1}{M} \int_A \pi(y) dy}{\frac{1}{M}} = \int_A \pi(y) dy.$$

$\square$

The acceptance probability is $\mathbb{P}(\text{accept}) = 1/M$, so the waiting time follows $\text{Geo}(1/M)$ with mean M . Given N i.i.d. proposals $Y_1, \dots, Y_N \sim g$ and $U_1, \dots, U_N \sim \text{Unif}(0, 1)$, with $r_i = \pi(Y_i)/(Mg(Y_i))$, the *RS estimator* of $\mathbb{E}_\pi[h]$ is

$$I_{\pi, N}^{\text{RS}}(h) = \frac{\sum_{i=1}^N \mathbf{1}\{U_i \leq r_i\} h(Y_i)}{\sum_{i=1}^N \mathbf{1}\{U_i \leq r_i\}}. \quad (2.1)$$

Remark 2.3. Rejected proposals are completely discarded. There is no mechanism for moving bad samples toward π ; one simply waits for a good draw. When M is large—as it inevitably is in high dimensions—this is very wasteful.

2.3 Importance Sampling

Known target. The key identity is

$$\mathbb{E}_\pi[h] = \int h(x) \pi(x) dx = \mathbb{E}_g \left[h(Y) \frac{\pi(Y)}{g(Y)} \right].$$

Setting $w_i = \pi(Y_i)/g(Y_i)$ and using $\mathbb{E}_g[w_i] = 1$, the *self-normalized IS estimator* is

$$I_{\pi,N}^{\text{IS}}(h) = \frac{\sum_{i=1}^N w_i h(Y_i)}{\sum_{i=1}^N w_i}. \quad (2.2)$$

Unknown normalization (AIS). In practice $\pi = \tilde{\pi}/Z$ with $\tilde{\pi} = e^{-V}$ computable but Z unknown. The unnormalized weights $\tilde{w}_i = \tilde{\pi}(Y_i)/g(Y_i)$ give the *auto-normalized IS (AIS) estimator*

$$I_{\pi,N}^{\text{AIS}}(h) = \frac{\sum_{i=1}^N \tilde{w}_i h(Y_i)}{\sum_{i=1}^N \tilde{w}_i}, \quad (2.3)$$

where Z cancels in the ratio.

Remark 2.4. Unlike RS, IS assigns every proposal a weight rather than accepting or rejecting it. A “bad” proposal (small π/g) contributes little but is not wasted. The cost is high weight variance when π and g differ significantly.

2.4 Asymptotic Theory

Both estimators are ratios of sample averages of i.i.d. variables. Their asymptotics follow from the multivariate CLT and the delta method for $(a, b) \mapsto a/b$.

Theorem 2.5 (CLT for RS and AIS). *Let $h \in L^2(\pi)$. As $N \rightarrow \infty$,*

$$\sqrt{N}(I_{\pi,N}^{\text{RS}}(h) - \mathbb{E}_\pi[h]) \xrightarrow{\mathcal{D}} \mathcal{N}(0, G_{\text{RS}}^2), \quad (2.4)$$

$$\sqrt{N}(I_{\pi,N}^{\text{AIS}}(h) - \mathbb{E}_\pi[h]) \xrightarrow{\mathcal{D}} \mathcal{N}(0, G_{\text{IS}}^2), \quad (2.5)$$

where

$$G_{\text{RS}}^2 = M \int_{\mathbb{R}^d} \pi(x) |h(x) - \mathbb{E}_\pi[h]|^2 dx, \quad (2.6)$$

$$G_{\text{IS}}^2 = \int_{\mathbb{R}^d} \frac{\pi(x)^2}{g(x)} |h(x) - \mathbb{E}_\pi[h]|^2 dx. \quad (2.7)$$

Remark 2.6. The factor M in G_{RS}^2 is the price of rejection: only a fraction $1/M$ of proposals contribute, inflating the variance by M .

Proposition 2.7 (Importance beats rejection). $G_{\text{IS}}^2 \leq G_{\text{RS}}^2$.

See Exercise 1(c) for the proof via Cauchy–Schwarz.

2.5 Non-Asymptotic Bound and the Role of χ^2

The CLT describes large- N fluctuations but does not say how large N must be. The following finite- N bound makes the dependence on N explicit and non-asymptotic.

Theorem 2.8 (Non-asymptotic MSE bound). For $\|h\|_\infty \leq 1$,

$$\mathbb{E} \left[(I_{\pi,N}^{\text{AIS}}(h) - \mathbb{E}_\pi[h])^2 \right] \leq \frac{4}{N} (\chi^2(\pi \| g) + 1), \quad (2.8)$$

where $\chi^2(\pi \| g) = \int \pi(x)^2 / g(x) dx - 1$.

To achieve $\text{MSE} \leq \varepsilon^2$, one needs $N \gtrsim (\chi^2(\pi \| g) + 1) / \varepsilon^2$. Since $G_{\text{IS}}^2 \leq 4(\chi^2(\pi \| g) + 1)$ for $\|h\|_\infty \leq 1$, the χ^2 -divergence controls both the asymptotic and non-asymptotic efficiency of AIS. See the next section for more definitions and discussions of χ^2 and KL.

2.6 The Curse of Dimensionality

Example 2.9 (Product measures). Let $\pi = \pi_1^{\otimes d}$, $g = g_1^{\otimes d}$, $M_1 = \sup_{x \in \mathbb{R}} \pi_1(x) / g_1(x)$, $c_1 = \chi^2(\pi_1 \| g_1)$. Then

$$\sup_{x \in \mathbb{R}^d} \frac{\pi(x)}{g(x)} = M_1^d, \quad \chi^2(\pi \| g) = (1 + c_1)^d - 1.$$

Both grow exponentially in d : RS needs M_1^d proposals per accepted sample, and AIS needs $N \sim (1 + c_1)^d / \varepsilon^2$ samples to achieve $\text{MSE} \leq \varepsilon^2$.

Example 2.10 (Shifted Gaussians). Let $\pi = \mathcal{N}(m_1, I_d)$, $g = \mathcal{N}(m_2, I_d)$. Completing the square gives

$$\chi^2(\pi \| g) = e^{\|m_1 - m_2\|^2} - 1, \quad D_{\text{KL}}(\pi \| g) = \frac{1}{2} \|m_1 - m_2\|^2.$$

For $m_1 - m_2 = (1, \dots, 1)^\top$ so $\|m_1 - m_2\|^2 = d$: AIS requires $N \sim e^d / \varepsilon^2$ samples, even though $D_{\text{KL}}(\pi \| g) = d/2$ grows only linearly.

The χ^2 -divergence, not KL, is the correct measure of IS difficulty: a proposal “close” to π in KL may have exponentially large χ^2 , rendering IS completely impractical.

The fundamental obstacle. Both RS and IS use a *static* proposal: independent draws from g are accepted or reweighted. There is no mechanism to move bad proposals toward π . This motivates *dynamical* approaches, where a distribution ρ_t evolves from an easy ρ_0 toward π via an ODE or SDE:

$$dX_t = b_t(X_t) dt \quad (\text{deterministic flow}), \quad (2.9)$$

$$dX_t = b_t(X_t) dt + \sigma_t dW_t \quad (\text{stochastic flow}). \quad (2.10)$$

The rest of the course constructs and analyzes such dynamics for high-dimensional sampling and generative modeling.

3 Probability Divergences

3.1 f -Divergences

Having seen that $\chi^2(\pi \| g)$ governs IS efficiency, we place it in the broader framework of f -divergences that can be used to measure dissimilarity between two measures.

Definition 3.1 (f -divergence). *Let $f : (0, \infty) \rightarrow \mathbb{R}$ be convex with $f(1) = 0$. The f -divergence between $P \ll Q$ with densities p, q on $\mathbb{R}^d$ is*

$$D_f(P \| Q) = \int_{\mathbb{R}^d} f\left(\frac{p(x)}{q(x)}\right) q(x) dx = \mathbb{E}_Q \left[f\left(\frac{p(Y)}{q(Y)}\right) \right]. \quad (3.1)$$

If $P \not\ll Q$, set $D_f(P \| Q) = +\infty$.

Proposition 3.2 (Basic properties). (i) Non-negativity: $D_f(P \| Q) \geq 0$.

(ii) Definiteness: if f is strictly convex at 1, then $D_f(P \| Q) = 0 \iff P = Q$.

Both follow from Jensen: $D_f(P \| Q) \geq f(\mathbb{E}_Q[p/q]) = f(1) = 0$.

3.2 Special Cases

$f(t)$	Name	Formula
$t \log t$	KL divergence $D_{\text{KL}}(P \ Q)$	$\int p \log \frac{p}{q} dx$
$-\log t$	Reverse KL $D_{\text{KL}}(Q \ P)$	$\int q \log \frac{q}{p} dx$
$(t-1)^2$	χ^2 -divergence $\chi^2(P \ Q)$	$\int \frac{(p-q)^2}{q} dx = \int \frac{p^2}{q} dx - 1$
$\frac{1}{2} t-1 $	Total variation $D_{\text{TV}}(P, Q)$	$\frac{1}{2} \int p - q dx$

3.3 Comparison and High-Dimensional Behavior

Proposition 3.3 (Ordering). *For any $P \ll Q$,*

$$4 D_{\text{TV}}(P, Q)^2 \leq D_{\text{KL}}(P \| Q) \leq \log(1 + \chi^2(P \| Q)) \leq \chi^2(P \| Q). \quad (3.2)$$

The first inequality is Pinsker's. The second follows from the log-sum inequality (Exercise 4(d)). The third uses $\log(1+t) \leq t$. In particular, χ^2 is strictly stronger than KL, which is stronger than TV.

For product measures $\pi = \pi_1^{\otimes d}$, $g = g_1^{\otimes d}$:

$$D_{\text{KL}}(\pi \| g) = d \cdot D_{\text{KL}}(\pi_1 \| g_1) \text{ (linear)}, \quad \chi^2(\pi \| g) = (1 + \chi^2(\pi_1 \| g_1))^d - 1 \text{ (exponential)}.$$

For shifted Gaussians with $\|m_1 - m_2\|^2 = d$:

$$D_{\text{TV}}(\pi, g) = O(1), \quad D_{\text{KL}}(\pi \| g) = \frac{d}{2}, \quad \chi^2(\pi \| g) = e^d - 1.$$

TV and KL give a benign picture; χ^2 reveals the true exponential difficulty for IS.

3.4 Preview: Density Evolution under ODE/SDE Dynamics

Returning to the dynamical approach (2.9)–(2.10), a key question is: how does the density ρ_t of X_t evolve?

Theorem 3.4 (Continuity and Fokker–Planck equations). *Let ρ_t denote the density of X_t .*

(i) **ODE** $dX_t = b_t(X_t) dt$: ρ_t satisfies the continuity equation

$$\partial_t \rho_t + \nabla \cdot (\rho_t b_t) = 0. \quad (3.3)$$

(ii) **SDE** $dX_t = b_t(X_t) dt + \sigma_t dW_t$: ρ_t satisfies the Fokker–Planck equation

$$\partial_t \rho_t + \nabla \cdot (\rho_t b_t) = \frac{1}{2} \sigma_t^2 \Delta \rho_t. \quad (3.4)$$

Proof sketch of (i). For $\phi \in C_c^\infty(\mathbb{R}^d)$,

$$\begin{aligned} \int \phi \partial_t \rho_t dx &= \frac{d}{dt} \mathbb{E}[\phi(X_t)] = \mathbb{E}[\nabla \phi(X_t) \cdot b_t(X_t)] \\ &= \int \nabla \phi \cdot b_t \rho_t dx = - \int \phi \nabla \cdot (\rho_t b_t) dx. \end{aligned}$$

Since ϕ is arbitrary, $\partial_t \rho_t = -\nabla \cdot (\rho_t b_t)$. □

Remark 3.5. Equation (3.3) expresses conservation of mass: ρ_t is transported by b_t with no creation or destruction. Equation (3.4) adds diffusion $\frac{\sigma_t^2}{2} \Delta \rho_t$ from the Brownian noise.

The rest of the course constructs b_t, σ_t so that $\rho_t \rightarrow \pi$ efficiently in high dimensions. We will discuss how these lead to dynamical algorithms and how these algorithms achieve accuracy in some f -divergence with certain complexity depending on dimensions.

Week 1 Exercises

1. (CLT for RS and AIS.) Prove Theorem 2.5 by the following steps.

(a) Write $I_{\pi,N}^{\text{RS}}(h) = A_N/B_N$ where

$$A_N = \frac{1}{N} \sum_{i=1}^N \mathbf{1}\{U_i \leq r_i\} h(Y_i), \quad B_N = \frac{1}{N} \sum_{i=1}^N \mathbf{1}\{U_i \leq r_i\}.$$

Verify the summands are i.i.d. under g . Compute $\mathbb{E}[A_N]$, $\mathbb{E}[B_N]$, and the 2×2 covariance matrix Σ of (A_N, B_N) . Apply the multivariate CLT and the delta method with $\varphi(a, b) = a/b$ to derive G_{RS}^2 in (2.6).

(b) Repeat for $I_{\pi,N}^{\text{AIS}}(h)$: identify the i.i.d. summands, compute their joint covariance, and apply the delta method to derive G_{IS}^2 in (2.7).

(c) Prove $G_{\text{IS}}^2 \leq G_{\text{RS}}^2$ using Cauchy–Schwarz. Characterize when equality holds.

2. (Non-asymptotic bound for AIS.) Prove Theorem 2.8: for $\|h\|_\infty \leq 1$,

$$\mathbb{E} \left[\left(I_{\pi,N}^{\text{AIS}}(h) - \mathbb{E}_\pi[h] \right)^2 \right] \leq \frac{4}{N} (\chi^2(\pi \| g) + 1).$$

Hint. Write $I_{\pi,N}^{\text{AIS}}(h) = \tilde{A}_N/\tilde{B}_N$ and note $I_{\pi,N}^{\text{AIS}}(h) - \mathbb{E}_\pi[h] = (\tilde{A}_N - \mathbb{E}_\pi[h] \cdot \tilde{B}_N)/\tilde{B}_N$. Argue $\tilde{B}_N > 0$ a.s., bound the numerator's second moment using $\|h\|_\infty \leq 1$, and show $\mathbb{E}_g[\tilde{w}_i^2]/Z^2 = \chi^2(\pi \| g) + 1$. See [18].

3. (χ^2 -divergence in high dimensions.)

(a) Let $\pi = \pi_1^{\otimes d}$ and $g = g_1^{\otimes d}$. Show $\chi^2(\pi \| g) = (1 + \chi^2(\pi_1 \| g_1))^d - 1$. *Hint.* Use $\int p_1^{\otimes d}(x)^2/g_1^{\otimes d}(x) dx = (\int p_1^2/g_1 dx_1)^d$. Conclude χ^2 is exponential in d while $D_{\text{KL}}(\pi \| g) = d \cdot D_{\text{KL}}(\pi_1 \| g_1)$ is linear.

(b) Let $\pi = \mathcal{N}(m_1, I_d)$ and $g = \mathcal{N}(m_2, I_d)$. By completing the square, show $\chi^2(\pi \| g) = e^{\|m_1 - m_2\|^2} - 1$. For $m_1 - m_2 = (1, \dots, 1)^\top$, how many AIS samples are needed to achieve $\text{MSE} \leq \varepsilon^2$?

(c) For the same Gaussian pair, compute $D_{\text{KL}}(\pi \| g)$ and $D_{\text{TV}}(\pi, g)$ as functions of $\delta = \|m_1 - m_2\|$. For $\delta^2 = d$, compare all three divergences as $d \rightarrow \infty$ and explain why $D_{\text{KL}}(\pi \| g) = O(d)$ is misleadingly optimistic about IS difficulty.

4. (f -divergence properties.)

(a) Prove $D_f(P \| Q) \geq 0$ using Jensen's inequality. State precisely where convexity and $f(1) = 0$ are used.

- (b) Prove that if f is strictly convex at 1, then $D_f(P \parallel Q) = 0 \implies P = Q$ as measures.
- (c) Verify each row of the table in Section 3 by substituting the given f into Theorem 3.1.
- (d) Prove $D_{\text{KL}}(P \parallel Q) \leq \log(1 + \chi^2(P \parallel Q))$. *Hint.* Write $p/q = 1 + (p - q)/q$ and use $\log t \leq t - 1$. Deduce $D_{\text{KL}}(P \parallel Q) \leq \chi^2(P \parallel Q)$.

5. (Optimal IS proposal.) Fix target π and h with $\mathbb{E}_\pi[h] = \mu_h$.

- (a) Show G_{IS}^2 is minimized over densities g with $\text{supp}(\pi) \subseteq \text{supp}(g)$ by $g^*(x) \propto \pi(x)|h(x) - \mu_h|$. *Hint.* Apply Cauchy–Schwarz: $\int \frac{(\pi|h - \mu_h|)^2}{g} dx \cdot \int g dx \geq (\int \pi|h - \mu_h| dx)^2$.
- (b) Compute $G_{\text{IS}}^2(g^*)$ in closed form.
- (c) Explain why g^* is not directly usable in practice.
- (d) When $h \geq 0$, find a proposal $g^\dagger$ achieving $G_{\text{IS}}^2(g^\dagger) = 0$. Why is $g^\dagger$ also impractical?

6. (Continuity and Fokker–Planck equations.)

- (a) Give a careful proof of Theorem 3.4(i) using a test function $\phi \in C_c^\infty(\mathbb{R}^d)$ and integration by parts. State which property each step uses.
- (b) Show that a stationary density ρ_∞ of (3.3) satisfies $\nabla \cdot (\rho_\infty b) = 0$. Can we find a nontrivial drift b such that $\pi \propto e^{-V}$ is stationary?
- (c) Prove (3.4) using Itô’s formula: $d\phi(X_t) = \nabla \phi \cdot b_t dt + \frac{\sigma_t^2}{2} \Delta \phi dt + \sigma_t \nabla \phi \cdot dW_t$.
- (d) Verify that for the overdamped Langevin SDE $dX_t = -\nabla V(X_t) dt + \sqrt{2} dW_t$, the stationary distribution of (3.4) is $\pi \propto e^{-V}$.

4 Langevin Dynamics

4.1 The Overdamped Langevin SDE

Recall the fundamental question: given query access to ∇V , how do we construct a drift b_t and noise level σ_t in (2.10) so that $\rho_t \rightarrow \pi \propto e^{-V}$?

Intuition from optimization. If we simply want to find the mode $x^* = \arg \min V$, we run gradient descent $\dot{x}_t = -\nabla V(x_t)$, which drives $x_t \rightarrow x^*$ exponentially fast under strong convexity. But sampling from π requires exploring the entire distribution, not just its peak. The *overdamped Langevin SDE* adds Brownian noise to gradient descent:

$$dX_t = -\nabla V(X_t) dt + \sqrt{2} dW_t. \quad (4.1)$$

The drift $-\nabla V$ pushes X_t toward regions of high probability, while the noise $\sqrt{2} dW_t$ prevents collapse and drives exploration. The factor $\sqrt{2}$ is chosen so that $\pi \propto e^{-V}$ is exactly stationary, as we now verify.

4.2 Fokker–Planck Equation and Stationarity

Derivation of the Fokker–Planck equation. Let ρ_t denote the density of X_t under (4.1). For any test function $\phi \in C_c^\infty(\mathbb{R}^d)$, we compute $\frac{d}{dt} \mathbb{E}[\phi(X_t)]$ in two ways.

On one hand, differentiating under the integral:

$$\frac{d}{dt} \mathbb{E}[\phi(X_t)] = \int \phi(x) \partial_t \rho_t(x) dx.$$

On the other hand, Itô's formula applied to (4.1) gives

$$d\phi(X_t) = \underbrace{\nabla \phi(X_t) \cdot (-\nabla V(X_t)) dt + \Delta \phi(X_t) dt}_{\text{bounded variation}} + \underbrace{\sqrt{2} \nabla \phi(X_t) \cdot dW_t}_{\text{martingale}}.$$

Taking expectations, the martingale term vanishes, so

$$\frac{d}{dt} \mathbb{E}[\phi(X_t)] = \int (-\nabla \phi \cdot \nabla V + \Delta \phi) \rho_t dx.$$

Integrating by parts (boundary terms vanish since $\phi \in C_c^\infty$):

$$\int (-\nabla \phi \cdot \nabla V) \rho_t dx = \int \phi \nabla \cdot (\rho_t \nabla V) dx, \quad \int \Delta \phi \rho_t dx = \int \phi \Delta \rho_t dx.$$

Since ϕ is arbitrary, equating the two expressions gives the *Fokker–Planck equation* for Langevin dynamics:

$$\boxed{\partial_t \rho_t = \nabla \cdot (\rho_t \nabla V) + \Delta \rho_t.} \quad (4.2)$$

Compact form via the score. Using the *score trick* $\Delta\rho = \nabla \cdot (\nabla\rho) = \nabla \cdot (\rho \nabla \log \rho)$, equation (10.12) becomes

$$\partial_t \rho_t = \nabla \cdot (\rho_t (\nabla V + \nabla \log \rho_t)) = \nabla \cdot \left(\rho_t \nabla \log \frac{\rho_t}{\pi} \right), \quad (4.3)$$

where the last step uses $\nabla V + \nabla \log \rho_t = -\nabla \log \pi + \nabla \log \rho_t = \nabla \log(\rho_t/\pi)$. This form immediately shows that $\rho_t = \pi$ is a stationary solution (the right-hand side vanishes), and that the dynamics transports mass along $-\nabla \log(\rho_t/\pi)$, pushing ρ_t toward π .

Proposition 4.1 (Stationarity of π). $\pi \propto e^{-V}$ is a stationary solution of (10.12). Equivalently, if $X_0 \sim \pi$ then $X_t \sim \pi$ for all $t \geq 0$.

Remark 4.2. The compact form (4.3) reveals that Langevin dynamics is a *gradient flow* of $D_{\text{KL}}(\rho_t \parallel \pi)$ in the space of probability measures equipped with the Wasserstein-2 metric (Jordan–Kinderlehrer–Otto, 1998) [12]. We will revisit this perspective when studying variational inference.

4.3 Mixing Time and the Wasserstein Metric

We now ask: starting from $\rho_0 \neq \pi$, how quickly does $\rho_t \rightarrow \pi$? To measure convergence we use the *Wasserstein-2 metric*, which is well-suited to the coupling arguments below and connects naturally to optimal transport.

Definition 4.3 (Wasserstein-2 metric). For probability measures μ, ν on $\mathbb{R}^d$,

$$W_2(\mu, \nu) = \left(\inf_{\gamma \in \Gamma(\mu, \nu)} \int \|x - y\|_2^2 d\gamma(x, y) \right)^{1/2} = \inf_{X \sim \mu, Y \sim \nu} (\mathbb{E} \|X - Y\|_2^2)^{1/2}, \quad (4.4)$$

where $\Gamma(\mu, \nu)$ denotes all couplings of μ and ν . The minimizing coupling is the optimal transport plan.

Example 4.4. $W_2(\delta_x, \delta_y) = \|x - y\|_2$. For shifted Gaussians, $W_2(\mathcal{N}(m_1, I_d), \mathcal{N}(m_2, I_d)) = \|m_1 - m_2\|_2$, growing like $\sqrt{d}$ for typical shifts, whereas $D_{\text{KL}}(\mathcal{N}(m_1, I_d) \parallel \mathcal{N}(m_2, I_d)) = \frac{1}{2} \|m_1 - m_2\|^2$ grows like d .

The *mixing time* to tolerance ε in metric D is

$$t_{\text{mix}}^D(\varepsilon; \rho_0) = \inf \{ t \geq 0 : D(\rho_t, \pi) \leq \varepsilon \}. \quad (4.5)$$

4.4 Convergence under Strong Convexity

Definition 4.5 (Strong convexity). $V \in C^2(\mathbb{R}^d)$ is α -strongly convex if $\nabla^2 V(x) \succeq \alpha I_d$ for all $x \in \mathbb{R}^d$.

Intuition from optimization. When V is α -strongly convex, the gradient flow ODE $\dot{x}_t = -\nabla V(x_t)$ satisfies

$$\frac{d}{dt} \|x_t - x^*\|_2^2 = 2(x_t - x^*) \cdot (-\nabla V(x_t)) \leq -2\alpha \|x_t - x^*\|_2^2,$$

using $(x - y) \cdot (\nabla V(x) - \nabla V(y)) \geq \alpha \|x - y\|_2^2$ and $\nabla V(x^*) = 0$, giving $\|x_t - x^*\|_2 \leq e^{-\alpha t} \|x_0 - x^*\|_2$. The Langevin SDE extends this contraction to the level of probability distributions via a *synchronous coupling*: run two copies of (4.1) with the same Brownian motion so that their noise cancels, reducing the problem to the deterministic ODE argument above.

Theorem 4.6 (W_2 contraction for Langevin). *Assume V is α -strongly convex. Let ρ_t be the law of X_t solving (4.1) with $X_0 \sim \rho_0$. Then*

$$W_2(\rho_t, \pi) \leq e^{-\alpha t} W_2(\rho_0, \pi). \quad (4.6)$$

Proof. Run two copies of (4.1) with the same W_t : $X_0 \sim \rho_0$ and $Y_0 \sim \pi$, so $Y_t \sim \pi$ for all t by stationarity. Since both processes share the same noise, the difference $Z_t = X_t - Y_t$ satisfies the *deterministic* ODE

$$\dot{Z}_t = -\nabla V(X_t) + \nabla V(Y_t).$$

Then $\frac{d}{dt} \|Z_t\|_2^2 = 2(X_t - Y_t) \cdot (-\nabla V(X_t) + \nabla V(Y_t)) \leq -2\alpha \|Z_t\|_2^2$ by strong convexity, so Gronwall gives $\mathbb{E}\|X_t - Y_t\|_2^2 \leq e^{-2\alpha t} \mathbb{E}\|X_0 - Y_0\|_2^2$. Optimizing over the initial coupling (X_0, Y_0) with $X_0 \sim \rho_0$, $Y_0 \sim \pi$ yields (4.6). $\square$

4.5 Initialization and Dimension Dependence

From (4.6), $W_2(\rho_t, \pi) \leq \varepsilon$ after $t \geq \frac{1}{\alpha} \log \frac{W_2(\rho_0, \pi)}{\varepsilon}$. It remains to bound $W_2(\rho_0, \pi)$ for a concrete initialization. For α -strongly convex V , a natural choice is $\rho_0 = \delta_{x^*}$.

Lemma 4.7 (Stein's identity). *For $X \sim \pi \propto e^{-V}$ and smooth $f : \mathbb{R}^d \rightarrow \mathbb{R}^d$,*

$$\mathbb{E}_\pi[-\nabla V(X) \cdot f(X)] = -\mathbb{E}_\pi[\nabla \cdot f(X)]. \quad (4.7)$$

Proof. Integration by parts: $\int (-\nabla V) \cdot f \pi dx = \int \nabla \log \pi \cdot f \pi dx = \int \nabla \pi \cdot f dx = -\int \pi \nabla \cdot f dx$. $\square$

Proposition 4.8 (Initial distance from the mode). *If V is α -strongly convex, then*

$$W_2^2(\delta_{x^*}, \pi) \leq \mathbb{E}_\pi[\|X - x^*\|_2^2] \leq \frac{d}{\alpha}.$$

Proof. The first inequality uses (x^*, X) with $X \sim \pi$ as a coupling. For the second, apply (4.7) with $f(x) = x - x^*$ (so $\nabla \cdot f = d$): $\mathbb{E}_\pi[-\nabla V(X) \cdot (X - x^*)] = -d$. Since $\nabla V(x^*) = 0$, strong convexity gives $(\nabla V(X) - \nabla V(x^*)) \cdot (X - x^*) \geq \alpha \|X - x^*\|_2^2$, so $\alpha \mathbb{E}_\pi \|X - x^*\|_2^2 \leq d$. $\square$

Corollary 4.9 (Mixing time from mode initialization). *Under α -strong convexity, starting from $\rho_0 = \delta_{x^*}$,*

$$t_{\text{mix}}^{W_2}(\varepsilon; \delta_{x^*}) \leq \frac{1}{\alpha} \log \frac{\sqrt{d/\alpha}}{\varepsilon}.$$

Remark 4.10 (Continuous time is not an algorithm). The mixing time in Theorem 4.9 is for the *continuous-time* SDE (4.1), which requires simulating the stochastic process exactly and is not directly implementable. In practice one must discretize (4.1) in time with step size $h > 0$, yielding the *Unadjusted Langevin Algorithm* (ULA):

$$X_{k+1} = X_k - h \nabla V(X_k) + \sqrt{2h} \xi_k, \quad \xi_k \stackrel{\text{i.i.d.}}{\sim} \mathcal{N}(0, I_d).$$

The true measure of algorithmic complexity is therefore the number of gradient evaluations k needed to reach $W_2(\text{law}(X_k), \pi) \leq \varepsilon$. This discrete-time analysis, which gives a practically meaningful complexity bound, is the subject of the next lecture. But just from the structure of the bound, the complexity will not scale with dimension exponentially, when the target distribution is strongly log-concave, i.e., V is strongly convex.

5 Langevin Monte Carlo

5.1 From Continuous Time to Algorithms

The continuous-time Langevin SDE (4.1) establishes that $\rho_t \rightarrow \pi$ exponentially fast, but it is not directly implementable. The standard discretization with step size $h > 0$ gives the *Unadjusted Langevin Algorithm* (ULA):

$$X_{(k+1)h} = X_{kh} - h \nabla V(X_{kh}) + \sqrt{2h} \xi_k, \quad \xi_k \stackrel{\text{i.i.d.}}{\sim} \mathcal{N}(0, I_d), \quad (5.1)$$

where $\sqrt{2h} \xi_k = \sqrt{2}(W_{(k+1)h} - W_{kh})$ is the discretized Brownian increment. Denote $\rho_{kh} = \text{law}(X_{kh})$.

Before analyzing ULA, we first revisit the simpler discrete-time setting of optimization, which provides the key structural insight.

5.2 Warm-up: Gradient Descent for Optimization

Under the assumption $\alpha I \preceq \nabla^2 V \preceq \beta I$, consider gradient descent (GD):

$$X_{(k+1)h} = X_{kh} - h \nabla V(X_{kh}). \quad (5.2)$$

To analyze convergence, write

$$\nabla V(X_{kh}) - \nabla V(x^*) = \underbrace{\left(\int_0^1 \nabla^2 V(sX_{kh} + (1-s)x^*) ds \right)}_{=: H_k} (X_{kh} - x^*),$$

where H_k satisfies $\alpha I \preceq H_k \preceq \beta I$. Since $\nabla V(x^*) = 0$, we have $\nabla V(X_{kh}) = H_k(X_{kh} - x^*)$, so

$$X_{(k+1)h} - x^* = (I - hH_k)(X_{kh} - x^*).$$

For $h \leq 1/\beta$, $\|I - hH_k\|_2 \leq 1 - h\alpha \leq e^{-h\alpha}$, giving

$$\|X_{kh} - x^*\|_2 \leq e^{-kh\alpha} \|X_0 - x^*\|_2. \quad (5.3)$$

To reach $\|X_{kh} - x^*\|_2 \leq \varepsilon$, it suffices to take $k \geq \frac{1}{h\alpha} \log \frac{\|X_0 - x^*\|_2}{\varepsilon}$, which for $h = 1/\beta$ gives $O\left(\frac{\beta}{\alpha} \log \frac{\|X_0 - x^*\|_2}{\varepsilon}\right)$ iterations. Here $\kappa = \beta/\alpha$ is called the condition number.

5.3 ULA: Main Theorem

We now state the main convergence result for ULA. Throughout, we assume $\alpha I \preceq \nabla^2 V \preceq \beta I$ with $\alpha, \beta > 0$, so $\pi \propto e^{-V}$ is *strongly log-concave*.

Theorem 5.1 (ULA convergence). *Under $\alpha I \preceq \nabla^2 V \preceq \beta I$, for step size $h \leq 1/\beta$, the ULA iterates satisfy*

$$W_2(\rho_{kh}, \pi) \leq e^{-kh\alpha} W_2(\rho_0, \pi) + O\left(\frac{\beta}{\alpha} \sqrt{hd}\right). \quad (5.4)$$

The bound (5.4) has two terms: an exponentially decaying *contraction* term (inherited from the continuous-time analysis), and a *discretization bias* term $O(\frac{\beta}{\alpha}\sqrt{hd})$ that is independent of k and persists even as $k \rightarrow \infty$. This confirms that ULA does not converge exactly to π , but to a neighborhood of size $O(\frac{\beta}{\alpha}\sqrt{hd})$ in W_2 .

5.4 Proof of Theorem 5.1

Proof. The proof proceeds by introducing a “true” (continuous-time) reference process Y_t , then splitting the error at each step into a contraction term and a discretization error via the triangle inequality.

Reference process. For $t \in [kh, (k+1)h]$, define Y_t as the solution to the Langevin SDE (4.1) on this interval, sharing the same Brownian motion as ULA:

$$dY_t = -\nabla V(Y_t) dt + \sqrt{2} dW_t, \quad kh \leq t \leq (k+1)h, \quad (5.5)$$

initialized so that (X_{kh}, Y_{kh}) is the *optimal coupling* of (ρ_{kh}, π) :

$$W_2(\rho_{kh}, \pi) = (\mathbb{E}\|X_{kh} - Y_{kh}\|_2^2)^{1/2}.$$

Since π is stationary for (5.5), $Y_t \sim \pi$ for all t , and in particular $Y_{(k+1)h} \sim \pi$.

Triangle inequality decomposition. Writing $Y_{(k+1)h}$ explicitly:

$$Y_{(k+1)h} = Y_{kh} - \int_{kh}^{(k+1)h} \nabla V(Y_t) dt + \sqrt{2}(W_{(k+1)h} - W_{kh}).$$

Since $X_{(k+1)h}$ and $Y_{(k+1)h}$ share the same Brownian motion,

$$\begin{aligned} W_2(\rho_{(k+1)h}, \pi) &\leq (\mathbb{E}\|X_{(k+1)h} - Y_{(k+1)h}\|_2^2)^{1/2} \\ &= \left(\mathbb{E} \left\| \underbrace{X_{kh} - h\nabla V(X_{kh}) - Y_{kh} + h\nabla V(Y_{kh})}_{\textcircled{1}} - \underbrace{\int_{kh}^{(k+1)h} (\nabla V(Y_t) - \nabla V(Y_{kh})) dt}_{\textcircled{2}} \right\|_2^2 \right)^{1/2}. \end{aligned}$$

By the triangle inequality:

$$W_2(\rho_{(k+1)h}, \pi) \leq \underbrace{(\mathbb{E}\|\textcircled{1}\|_2^2)^{1/2}}_{\text{GD-like term}} + \underbrace{(\mathbb{E}\|\textcircled{2}\|_2^2)^{1/2}}_{\text{discretization error}}. \quad (5.6)$$

Term $\textcircled{1}$: GD-like contraction. By the same linearization as in (5.3), $\textcircled{1} = (I - hH_k)(X_{kh} - Y_{kh})$ where $H_k = \int_0^1 \nabla^2 V(sX_{kh} + (1-s)Y_{kh}) ds$ satisfies $\alpha I \preceq H_k \preceq \beta I$. For $h \leq 1/\beta$, $\|I - hH_k\|_2 \leq e^{-h\alpha}$, so

$$(\mathbb{E}\|\textcircled{1}\|_2^2)^{1/2} \leq e^{-h\alpha} W_2(\rho_{kh}, \pi). \quad (5.7)$$

Term ②: discretization error. We first bound the expected squared gradient norm.

Lemma 5.2 (Expected squared gradient norm). *For $Y \sim \pi \propto e^{-V}$ with $\alpha I \preceq \nabla^2 V \preceq \beta I$,*

$$\mathbb{E}_\pi \|\nabla V(Y)\|_2^2 \leq \beta d.$$

Proof. Apply Stein's identity (4.7) with $f(y) = \nabla V(y)$ (so $\nabla \cdot f = \Delta V$):

$$-\mathbb{E}_\pi[\nabla V(Y) \cdot \nabla V(Y)] = -\mathbb{E}_\pi[\Delta V(Y)],$$

giving $\mathbb{E}_\pi \|\nabla V(Y)\|_2^2 = \mathbb{E}_\pi[\Delta V(Y)]$. Since $\nabla^2 V \preceq \beta I$, $\Delta V(y) = \text{tr}(\nabla^2 V(y)) \leq \beta d$, so $\mathbb{E}_\pi \|\nabla V(Y)\|_2^2 \leq \beta d$. $\square$

Lemma 5.3 (Discretization error). *Under $\alpha I \preceq \nabla^2 V \preceq \beta I$ and $h \leq 1/\beta$,*

$$(\mathbb{E} \|\textcircled{2}\|_2^2)^{1/2} = O\left(\beta \sqrt{h^3 d}\right).$$

Proof. We have $\textcircled{2} = \int_{kh}^{(k+1)h} (\nabla V(Y_t) - \nabla V(Y_{kh})) dt$. By Cauchy-Schwarz in time and $\|\nabla^2 V\|_2 \leq \beta$:

$$\mathbb{E} \|\textcircled{2}\|_2^2 \leq h \beta^2 \int_{kh}^{(k+1)h} \mathbb{E} \|Y_t - Y_{kh}\|_2^2 dt.$$

Writing $Y_t - Y_{kh} = -\int_{kh}^t \nabla V(Y_s) ds + \sqrt{2}(W_t - W_{kh})$ and applying $(a+b)^2 \leq 2a^2 + 2b^2$:

$$\mathbb{E} \|Y_t - Y_{kh}\|_2^2 \leq 2\mathbb{E} \left\| \int_{kh}^t \nabla V(Y_s) ds \right\|_2^2 + 4(t - kh)d.$$

By Cauchy-Schwarz in time and Theorem 5.2: $\mathbb{E} \left\| \int_{kh}^t \nabla V(Y_s) ds \right\|_2^2 \leq h^2 \beta d$. Therefore $\mathbb{E} \|Y_t - Y_{kh}\|_2^2 \leq 6hd$ for $h \leq 1/\beta$, giving $\mathbb{E} \|\textcircled{2}\|_2^2 \leq 6h^3 \beta^2 d$. $\square$

Combining the bounds. Substituting (5.7) and Lemma 5.3 into (5.6) and unrolling over k steps, using $1 - e^{-h\alpha} \geq h\alpha/2$ to sum the geometric series:

$$W_2(\rho_{kh}, \pi) \leq e^{-kh\alpha} W_2(\rho_0, \pi) + O\left(\frac{\beta}{\alpha} \sqrt{hd}\right). \quad (5.8)$$

$\square$

5.5 Algorithmic Complexity

To reach $W_2(\rho_{kh}, \pi) \leq \varepsilon$, it suffices to require simultaneously that the contraction term satisfies $e^{-kh\alpha} W_2(\rho_0, \pi) \leq \varepsilon/2$ and the bias term satisfies $O(\frac{\beta}{\alpha} \sqrt{hd}) \leq \varepsilon/2$.

Corollary 5.4 (Iteration complexity of ULA). *Under $\alpha I \preceq \nabla^2 V \preceq \beta I$, initialize at $\rho_0 = \delta_{x^*}$ so that $W_2(\rho_0, \pi) \leq \sqrt{d/\alpha}$. Choose step size $h \asymp \alpha^2 \varepsilon^2 / (\beta^2 d)$. Then to achieve $W_2(\rho_{kh}, \pi) \leq \varepsilon$, the required number of gradient evaluations is*

$$k = O\left(\frac{\kappa^2 d}{\varepsilon^2} \log \frac{\sqrt{d/\alpha}}{\varepsilon}\right), \quad (5.9)$$

where $\kappa = \beta/\alpha$ is the condition number. For fixed ε and α , the dimension dependence is $O(d \log d)$.

Proof. With $h \asymp \alpha^2 \varepsilon^2 / (\beta^2 d)$, the bias satisfies $\frac{\beta}{\alpha} \sqrt{hd} \asymp \varepsilon$. The contraction condition $e^{-kh\alpha} W_2(\rho_0, \pi) \leq \varepsilon/2$ then requires $k \geq \frac{1}{h\alpha} \log \frac{W_2(\rho_0, \pi)}{\varepsilon} \asymp \frac{\kappa^2 d}{\varepsilon^2} \log \frac{\sqrt{d/\alpha}}{\varepsilon}$. $\square$

Remark 5.5 (Dimensional scaling and references). The $O(d \log d)$ iteration complexity is a significant improvement over the exponential complexity of IS and RS. Other results measuring convergence in KL or χ^2 are given in Chewi's book [7]. The dimensional dependence can be improved to $\tilde{O}(\sqrt{d})$ by more refined analysis and assumptions; see Li–Zha–Tao (2022) [13]. We also note a dimension-free result if the quantity of interest only depends on low-dimensional coordinates [5].

Week 2 Exercises

1. (Gaussian target: explicit computations.) Throughout, let $\pi = \mathcal{N}(0, \Sigma)$ with $\Sigma \succ 0$, so $V(x) = \frac{1}{2}x^\top \Sigma^{-1}x$ and $\nabla V(x) = \Sigma^{-1}x$.

- (a) (*Continuous-time Langevin.*) Show that the Langevin SDE (4.1) becomes the Ornstein–Uhlenbeck (OU) process $dX_t = -\Sigma^{-1}X_t dt + \sqrt{2} dW_t$. If $X_0 \sim \mathcal{N}(m_0, S_0)$, show that $X_t \sim \mathcal{N}(m_t, S_t)$ for all $t \geq 0$, where

$$m_t = e^{-\Sigma^{-1}t}m_0, \quad S_t = e^{-\Sigma^{-1}t}S_0 e^{-\Sigma^{-1}t} + \Sigma(I - e^{-2\Sigma^{-1}t}).$$

Verify that $m_t \rightarrow 0$ and $S_t \rightarrow \Sigma$ as $t \rightarrow \infty$.

- (b) (*Explicit W_2 convergence.*) For two Gaussians, $W_2^2(\mathcal{N}(\mu_1, \Sigma_1), \mathcal{N}(\mu_2, \Sigma_2)) = \|\mu_1 - \mu_2\|_2^2 + \|\Sigma_1^{1/2} - \Sigma_2^{1/2}\|_F^2$. Work in the eigenbasis of $\Sigma = \text{diag}(\lambda_1, \dots, \lambda_d)$ with $0 < \lambda_1 \leq \dots \leq \lambda_d$, and assume S_0 is diagonal with entries $s_1, \dots, s_d > 0$.

- (i) Show $\|m_t\|_2^2 = \sum_{i=1}^d e^{-2t/\lambda_i} m_{0,i}^2$, decaying at rate $2/\lambda_d = 2\lambda_{\min}(\Sigma^{-1})$.
(ii) Show $S_{t,ii} = e^{-2t/\lambda_i} s_i + \lambda_i(1 - e^{-2t/\lambda_i})$, and for large t ,

$$\sqrt{S_{t,ii}} - \sqrt{\lambda_i} \approx \frac{s_i - \lambda_i}{2\sqrt{\lambda_i}} e^{-2t/\lambda_i},$$

so $\|S_t^{1/2} - \Sigma^{1/2}\|_F^2$ decays at rate $4/\lambda_d$.

- (iii) Conclude $W_2(\rho_t, \pi) \leq C e^{-\alpha t}$ where $\alpha = 1/\lambda_d = \lambda_{\min}(\Sigma^{-1})$, confirming Theorem 4.6.

Remark: in the non-diagonal case the dominant decay rate $e^{-\alpha t}$ persists, but the formula for $\|S_t^{1/2} - \Sigma^{1/2}\|_F^2$ is more involved.

- (c) (*ULA stationary distribution.*) Specialize to $\Sigma = I_d$ ($\alpha = \beta = 1$). For ULA with step size $h > 0$, $X_{(k+1)h} = (1-h)X_{kh} + \sqrt{2h}\xi_k$, show that $X_{kh} \sim \mathcal{N}(0, S_{kh})$ with $S_{kh} \rightarrow S_\infty = \frac{2}{2-h}I_d$ as $k \rightarrow \infty$. Hence ULA has a biased stationary distribution $\pi_h = \mathcal{N}\left(0, \frac{2}{2-h}I_d\right) \neq \pi$.

- (d) (*Bias and optimal step size.*) Show that

$$W_2(\pi_h, \pi) = \sqrt{d} \left| \sqrt{\frac{2}{2-h}} - 1 \right| \asymp \frac{\sqrt{d}}{4} h \quad \text{as } h \rightarrow 0.$$

Conclude that $h \asymp \varepsilon/\sqrt{d}$ suffices to achieve bias $\leq \varepsilon$. Compare to the general bound $h \asymp \varepsilon^2/(\beta^2 d)$ from Theorem 5.4 (with $\beta = 1$) and explain the improvement.

2. (Convergence under convexity without strong convexity.) Assume $V \in C^2(\mathbb{R}^d)$ is convex ($\nabla^2 V \succeq 0$) with minimizer $x^* = \arg \min V$, but not α -strongly convex. Recall the gradient inequality

$$V(y) \geq V(x) + \nabla V(x) \cdot (y - x) \quad \text{for all } x, y \in \mathbb{R}^d. \quad (5.10)$$

- (a) (*Gradient flow: $O(1/T)$ rate.*) Consider $\dot{x}_t = -\nabla V(x_t)$.
- (i) Show $\frac{d}{dt}V(x_t) = -\|\nabla V(x_t)\|_2^2 \leq 0$.
 - (ii) Using (5.10) with $y = x^*$, show $V(x_t) - V(x^*) \leq -\frac{1}{2} \frac{d}{dt} \|x_t - x^*\|_2^2$.
 - (iii) Integrate and use monotonicity of $V(x_t)$ to deduce

$$V(x_T) - V(x^*) \leq \frac{\|x_0 - x^*\|_2^2}{2T}.$$

- (b) (*Discrete GD: $O(1/K)$ rate.*) Consider GD (5.2) with $h \leq 1/\beta$.
- (i) Using the descent lemma, show $V(X_{(k+1)h}) - V(X_{kh}) \leq -\frac{h}{2} \|\nabla V(X_{kh})\|_2^2$, so $V(X_{kh})$ is monotone decreasing.
 - (ii) Using (5.10), show $V(X_{kh}) - V(x^*) \leq \frac{1}{2h} (\|X_{kh} - x^*\|_2^2 - \|X_{(k+1)h} - x^*\|_2^2)$. Sum over $k = 0, \dots, K-1$, use telescoping and monotone decrease of $V(X_{kh})$ to obtain

$$V(X_{Kh}) - V(x^*) \leq \frac{\|X_0 - x^*\|_2^2}{2Kh}.$$

This $O(1/K)$ bound holds for the last iterate X_{Kh} .

- (c) (*Extension to the Langevin setting.*) Optional. This relies on the interpretation that Langevin can be viewed as a gradient flow of KL under the Wasserstein geometry which we haven't discussed in the lecture.
- (i) (*KL decay identity.*) Let $\mathcal{I}(\rho \parallel \pi) = \int \|\nabla \log(\rho/\pi)\|_2^2 \rho dx$ be the relative Fisher information. Using (4.3), prove

$$\frac{d}{dt} D_{\text{KL}}(\rho_t \parallel \pi) = -\mathcal{I}(\rho_t \parallel \pi).$$

This is the exact analogue of $\frac{d}{dt}V(x_t) = -\|\nabla V(x_t)\|_2^2$ from (a)(i).

- (ii) (*$O(1/T)$ convergence via the HWI inequality.*) Under convexity of V , the HWI inequality (Otto–Villani) states $D_{\text{KL}}(\rho \parallel \pi) \leq W_2(\rho, \pi) \sqrt{\mathcal{I}(\rho \parallel \pi)}$. Use the following steps to deduce $D_{\text{KL}}(\rho_T \parallel \pi) \leq W_2(\rho_0, \pi)^2/T$:
 - Use the synchronous coupling with $\alpha = 0$ to show $W_2(\rho_t, \pi) \leq W_2(\rho_0, \pi) =: C$ for all $t \geq 0$.
 - Let $f(t) = D_{\text{KL}}(\rho_t \parallel \pi)$. Apply HWI to get $f(t) \leq C \sqrt{\mathcal{I}(\rho_t \parallel \pi)}$, hence $\mathcal{I}(\rho_t \parallel \pi) \geq f(t)^2/C^2$ (squaring and rearranging). Combined with (i): $-f'(t) \geq f(t)^2/C^2$.
 - Rewrite as $\frac{d}{dt}(1/f(t)) \geq 1/C^2$, integrate from 0 to T , and conclude $D_{\text{KL}}(\rho_T \parallel \pi) \leq W_2(\rho_0, \pi)^2/T$.

Compare with (a)(iii): the arguments are structurally identical, with $(V - V^*, \|\nabla V\|^2, \|x - x^*\|)$ replaced by $(D_{\text{KL}}(\rho \parallel \pi), \mathcal{I}, W_2)$.

6 Score-Based Diffusion Models

6.1 Motivation and Setting

The setting of this section differs from the previous ones. We do *not* have access to a model V for the target π , but instead have data $X^{(1)}, \dots, X^{(N)} \sim \pi$ and wish to generate new samples from π . This is the *generative modeling* problem.

Popular approaches all try to find a transport from a simple reference $\gamma = \mathcal{N}(0, I_d)$ to the target π : variational autoencoders (VAE), generative adversarial networks (GAN), normalizing flows (NF), and diffusion models (DM). This section focuses on score based diffusion models, which approach this through a learned stochastic dynamics involving scores.

Recall from Section 4 that Langevin dynamics can be written as

$$dX_t = \nabla \log \pi(X_t) dt + \sqrt{2} dW_t,$$

so if we can learn *the score* $\nabla \log \pi$ from data, we can sample by running Langevin.

6.2 Score Matching

Warm-up: learning the score from data. Given a neural network $s(\cdot, \theta)$, we wish to find θ so that $s(x, \theta) \approx \nabla \log \pi(x)$. The natural loss is

$$L(\theta) = \int \|s(x, \theta) - \nabla \log \pi(x)\|_2^2 \pi(x) dx,$$

but this requires knowing $\nabla \log \pi$, which is unavailable. Integration by parts transforms this into a tractable form.

Proposition 6.1 (Score matching identity). *Up to a constant C independent of θ ,*

$$L(\theta) = \mathbb{E}_\pi [\|s(X, \theta)\|_2^2 + 2 \nabla \cdot s(X, \theta)] + C. \quad (6.1)$$

Proof. Expand $\|s - \nabla \log \pi\|^2 = \|s\|^2 - 2s \cdot \nabla \log \pi + \|\nabla \log \pi\|^2$. The last term is constant in θ . For the cross term,

$$\begin{aligned} \int -2s(x, \theta) \cdot \nabla \log \pi(x) \pi(x) dx &= -2 \int s(x, \theta) \cdot \nabla \pi(x) dx \\ &= 2 \int (\nabla \cdot s(x, \theta)) \pi(x) dx, \end{aligned}$$

where the last step is integration by parts (boundary terms vanish). Adding $\mathbb{E}_\pi[\|s\|^2]$ gives (6.1). $\square$

The loss (6.1) depends only on expectations under π , so it can be empiricalized using data and minimized with stochastic optimization algorithms such as SGD/ADAM.

Issues. Plain score matching has two significant drawbacks:

- ① *The divergence term $\nabla \cdot s(X, \theta)$ is expensive to evaluate* for high-dimensional neural networks. (Variants such as sliced score matching reduce this cost.)
- ② *The loss only guarantees $L^2(\pi)$ -accuracy:* since it averages against π , the score estimate may be inaccurate in low-probability regions where $\pi(x)$ is small. This is problematic.

Moreover, even with the exact score, running Langevin to sample from a multi-modal π ($\pi \propto e^{-V}$ with non-convex V) can be very slow: the chain must traverse low-probability barriers between modes.

6.3 Multiscale Approach and the Noising Process

A simple trick: add noise. To address issue ②, replace π by the smoothed target $\pi * \mathcal{N}(0, \sigma^2 I)$. Adding noise to data yields samples from this smoothed distribution. One then:

- estimates the score of $\pi * \mathcal{N}(0, \sigma^2 I)$ first (more well-behaved globally, as the low-probability region is smoothed and becomes less low-probability);
- runs Langevin with the learned score to get approximate samples, then refines using the original score. One can imagine the convergence will be faster using this annealing strategy.

This multiscale idea is the intuitive foundation of score based diffusion models [21, 22].

The OU noising process (Variance Preserving design). Rather than a single noise level, we consider a continuous-time noising process. The standard choice is the *Ornstein-Uhlenbeck* (OU) process:

$$dX_t = -X_t dt + \sqrt{2} dW_t, \quad X_0 \sim \pi, \quad (6.2)$$

which is Langevin dynamics (4.1) with $V(x) = \frac{1}{2}\|x\|^2$, so $\mathcal{N}(0, I_d)$ is stationary. As $t \rightarrow \infty$, the distribution of X_t converges to $\mathcal{N}(0, I_d)$ regardless of X_0 .

Proposition 6.2 (Marginal of the OU process). *Let $X_0 \sim \pi$ and $Z \sim \mathcal{N}(0, I_d)$ be independent. Then*

$$X_t \stackrel{d}{=} e^{-t} X_0 + \sqrt{1 - e^{-2t}} Z \sim \mathcal{N}(e^{-t} X_0, (1 - e^{-2t}) I_d) \text{ given } X_0. \quad (6.3)$$

In particular, the conditional distribution is $p_t(\cdot | x_0) = \mathcal{N}(e^{-t} x_0, (1 - e^{-2t}) I_d)$.

Proof. Apply Itô's formula to $e^t X_t$:

$$d(e^t X_t) = e^t X_t dt + e^t dX_t = e^t X_t dt + e^t (-X_t dt + \sqrt{2} dW_t) = \sqrt{2} e^t dW_t.$$

Integrating from 0 to t : $e^t X_t - X_0 = \int_0^t \sqrt{2} e^s dW_s \sim \mathcal{N}(0, 2 \int_0^t e^{2s} ds) = \mathcal{N}(0, e^{2t} - 1)$.

Therefore $X_t \stackrel{d}{=} e^{-t} X_0 + \sqrt{1 - e^{-2t}} Z$. $\square$

6.4 Time Reversal and the Reverse SDE

The key idea. If we can *reverse* the noising process (6.2), we can generate samples using it. The following theorem identifies the reverse dynamics.

Theorem 6.3 (Time reversal [3], Anderson 1982). *Let X_t solve (6.2) for $t \in [0, T]$ with $X_0 \sim \pi$, and let p_t denote the density of X_t . Define the time-reversal $Y_t := X_{T-t}$ for $t \in [0, T]$. Then Y_t solves the reverse SDE*

$$dY_t = (Y_t + 2\nabla \log p_{T-t}(Y_t)) dt + \sqrt{2} dW_t, \quad (6.4)$$

where W_t is a Brownian motion.

Heuristic demonstration via Fokker–Planck. Let $q_t(x) = p_{T-t}(x)$ be the density of Y_t . The OU process (6.2) satisfies the Fokker–Planck equation

$$\partial_t p_t = \nabla \cdot (p_t x) + \Delta p_t.$$

Since $\partial_t q_t(x) = -\partial_t p_{T-t}(x)$, we have

$$\partial_t q_t = -\nabla \cdot (q_t x) - \Delta q_t.$$

Using the score trick $\Delta q = \nabla \cdot (q \nabla \log q)$:

$$\begin{aligned} \partial_t q_t &= -\nabla \cdot (q_t x) - \nabla \cdot (q_t \nabla \log q_t) \\ &= -\nabla \cdot (q_t(x + 2\nabla \log q_t)) + \Delta q_t. \end{aligned}$$

This leads to

$$\partial_t q_t + \nabla \cdot (q_t(x + 2\nabla \log q_t)) = \Delta q_t,$$

which is the Fokker–Planck equation for the SDE (6.4) (with $q_t = p_{T-t}$). $\square$

Practice: the approximate reverse process. In practice the score $\nabla \log p_{T-t}$ is unknown and replaced by a learned approximation $\hat{S}_{T-t}(\cdot, \theta)$. The *approximate reverse process* is

$$d\hat{Y}_t = (\hat{Y}_t + 2\hat{S}_{T-t}(\hat{Y}_t, \theta)) dt + \sqrt{2} dW_t, \quad \hat{Y}_0 \sim \mathcal{N}(0, I_d). \quad (6.5)$$

Let $\hat{q}_t$ denote the law of $\hat{Y}_t$. The goal is $\hat{q}_T \approx q_T = p_0 = \pi$, which holds (intuitively) when three conditions are met:

- (1) T is large enough so that $p_T \approx \mathcal{N}(0, I_d)$;
- (2) $\hat{S}_t \approx \nabla \log p_t$ accurately enough;
- (3) the time discretization of (6.5) is fine enough.

6.5 Denoising Score Matching

It remains to learn $\nabla \log p_t$ tractably from data.

Marginalizing the score. Since $p_t(x) = \int p_t(x | x_0) p_0(x_0) dx_0$, differentiating gives

$$\nabla_x p_t(x) = \int \nabla_x p_t(x | x_0) p_0(x_0) dx_0.$$

Dividing by $p_t(x)$ and using the Bayes formula for conditional distributions:

$$\nabla \log p_t(x) = \mathbb{E}[\nabla_{X_t} \log p_t(X_t | X_0) | X_t = x]. \quad (6.6)$$

Since $p_t(\cdot | x_0) = \mathcal{N}(e^{-t}x_0, (1 - e^{-2t})I_d)$, the conditional score is explicit:

$$\nabla_x \log p_t(x | x_0) = -\frac{x - e^{-t}x_0}{1 - e^{-2t}}. \quad (6.7)$$

In particular,

$$\nabla_x \log p_t(x) = \frac{e^{-t}\mathbb{E}[X_0 | X_t = x] - x}{1 - e^{-2t}}. \quad (6.8)$$

Here $\mathbb{E}[X_0 | X_t]$ is the optimal estimator (denoiser) of the clean data X_0 given the noisy observation X_t . Hence *learning the score is equivalent to learning the denoiser* $\mathbb{E}[X_0 | X_t]$. This above is also referred to as one form of Tweedie's formula.

Denoising score matching loss. By (6.6), the marginal score is the best $L^2(p_t)$ -estimator given X_t . This motivates learning $S_t(x, \theta) \approx \nabla \log p_t(x)$ by minimizing the *denoising score matching* (DSM) loss:

$$L(\theta) = \int_0^T \mathbb{E}_{X_0 \sim p_0, X_t | X_0} \left[\|S_t(X_t, \theta) - \nabla_{X_t} \log p_t(X_t | X_0)\|_2^2 \right] dt. \quad (6.9)$$

By (6.6), $L(\theta)$ equals the original score matching objective integrated over time:

$$\int_0^T \mathbb{E}_{X_t \sim p_t} \left[\|S_t(X_t, \theta) - \nabla_{X_t} \log p_t(X_t)\|_2^2 \right] dt, \quad (6.10)$$

up to a constant that is independent of θ .

Substituting the explicit formula (6.7) into (6.9):

$$L(\theta) = \int_0^T \mathbb{E} \left[\left\| S_t(X_t, \theta) - \frac{e^{-t}X_0 - X_t}{1 - e^{-2t}} \right\|_2^2 \right] dt. \quad (6.11)$$

The loss (6.11) is fully tractable: $X_0 \sim p_0$ is sampled from data; $X_t | X_0 \sim \mathcal{N}(e^{-t}X_0, (1 - e^{-2t})I_d)$ is sampled via (6.3); and $\nabla_x \log p_t(X_t | X_0)$ is given by (6.7). There is no divergence term and no intractable expectation. In practice one samples a batch $X_0 \sim p_0$, a random time $t \in [0, T]$, and $X_t | X_0$, evaluates the empirical loss, and performs a step of SGD/ADAM. One often includes a time-dependent weight $w(t)$ to balance the loss at different noise levels.

6.6 Error Analysis

We now quantify how well $\hat{q}_T$ approximates $q_T = p_0 = \pi$ in terms of the score approximation error. Recall that $q_t = p_{T-t}$ denotes the law of the true time-reversal Y_t , and $\hat{q}_t$ denotes the law of the approximate process $\hat{Y}_t$ from (6.5). The drifts are

$$b_t(y) = y + 2\nabla \log p_{T-t}(y), \quad \hat{b}_t(y) = y + 2S_{T-t}(y, \theta), \quad (6.12)$$

so $q_0 = p_T$, $\hat{q}_0 = \mathcal{N}(0, I_d)$, $q_T = p_0 = \pi$, and $\hat{q}_T \approx \pi$ is the goal.

Proposition 6.4 (W_2 error bound). *Assume $\hat{b}_t$ is Lipschitz in x with constant $\hat{L}$, uniformly in t . Then*

$$W_2(q_T, \hat{q}_T)^2 \leq e^{(2\hat{L}+1)T} \left(W_2(q_0, \hat{q}_0)^2 + \int_0^T \mathbb{E}_{q_t} [\|b_t(Y_t) - \hat{b}_t(Y_t)\|_2^2] dt \right). \quad (6.13)$$

The integral on the right is exactly the score matching loss.

Proof. Couple (6.4) and (6.5) with the same W_t . Then

$$\frac{d}{dt} \|Y_t - \hat{Y}_t\|_2^2 = 2(Y_t - \hat{Y}_t) \cdot (b_t(Y_t) - \hat{b}_t(\hat{Y}_t)).$$

Add and subtract $\hat{b}_t(Y_t)$:

$$= 2(Y_t - \hat{Y}_t) \cdot (b_t(Y_t) - \hat{b}_t(Y_t)) + 2(Y_t - \hat{Y}_t) \cdot (\hat{b}_t(Y_t) - \hat{b}_t(\hat{Y}_t)).$$

For the first term, Young's inequality $2ab \leq a^2 + b^2$ gives

$$2(Y_t - \hat{Y}_t) \cdot (b_t(Y_t) - \hat{b}_t(Y_t)) \leq \|Y_t - \hat{Y}_t\|_2^2 + \|b_t(Y_t) - \hat{b}_t(Y_t)\|_2^2.$$

For the second, the Lipschitz condition on $\hat{b}_t$ gives

$$2(Y_t - \hat{Y}_t) \cdot (\hat{b}_t(Y_t) - \hat{b}_t(\hat{Y}_t)) \leq 2\hat{L}\|Y_t - \hat{Y}_t\|_2^2.$$

Combining:

$$\frac{d}{dt} \|Y_t - \hat{Y}_t\|_2^2 \leq (2\hat{L} + 1)\|Y_t - \hat{Y}_t\|_2^2 + \|b_t(Y_t) - \hat{b}_t(Y_t)\|_2^2.$$

Taking expectations and applying Gronwall's inequality:

$$\mathbb{E}\|Y_T - \hat{Y}_T\|_2^2 \leq e^{(2\hat{L}+1)T} \left(\mathbb{E}\|Y_0 - \hat{Y}_0\|_2^2 + \int_0^T \mathbb{E}[\|b_t(Y_t) - \hat{b}_t(Y_t)\|_2^2] dt \right).$$

Since $\mathbb{E}\|Y_T - \hat{Y}_T\|_2^2 \geq W_2(q_T, \hat{q}_T)^2$ and minimizing $\mathbb{E}\|Y_0 - \hat{Y}_0\|_2^2$ over all couplings of $(q_0, \hat{q}_0)$ gives $W_2(q_0, \hat{q}_0)^2$, we obtain (6.13). $\square$

The above W_2 upper bound using coupling arguments leads to exponential dependence on T .

Proposition 6.5 (KL error bound). *With the same setup,*

$$D_{\text{KL}}(q_T \parallel \hat{q}_T) \leq D_{\text{KL}}(q_0 \parallel \hat{q}_0) + \int_0^T \frac{1}{4} \mathbb{E}_{q_t} [\|b_t(Y_t) - \hat{b}_t(Y_t)\|_2^2] dt. \quad (6.14)$$

Proof. This can be easily derived using path-KL and Girsanov formula:

$$D_{\text{KL}}(q_T \parallel \hat{q}_T) \leq D_{\text{KL}}(q_{[0,T]} \parallel \hat{q}_{[0,T]}) = D_{\text{KL}}(q_0 \parallel \hat{q}_0) + \int_0^T \frac{1}{4} \mathbb{E}_{q_t} [\|b_t(Y_t) - \hat{b}_t(Y_t)\|_2^2] dt.$$

We also provide an explicit proof based on PDE calculations. Both q_t and $\hat{q}_t$ satisfy Fokker–Planck equations

$$\partial_t q_t + \nabla \cdot (q_t b_t) = \Delta q_t, \quad \partial_t \hat{q}_t + \nabla \cdot (\hat{q}_t \hat{b}_t) = \Delta \hat{q}_t.$$

We differentiate $D_{\text{KL}}(q_t \parallel \hat{q}_t) = \int q_t \log(q_t/\hat{q}_t) dx$:

$$\frac{d}{dt} D_{\text{KL}}(q_t \parallel \hat{q}_t) = \underbrace{\int \partial_t q_t \log \frac{q_t}{\hat{q}_t} dx}_{(A)} - \underbrace{\int \frac{q_t}{\hat{q}_t} \partial_t \hat{q}_t dx}_{(B)}, \quad (6.15)$$

where the middle term $\int \partial_t q_t dx = 0$ was dropped.

Term (A). Substituting the FP equation for q_t and integrating by parts:

$$\begin{aligned} (A) &= \int \left(-\nabla \cdot (q_t b_t) + \Delta q_t \right) \log \frac{q_t}{\hat{q}_t} dx \\ &= \int q_t b_t \cdot \nabla \log \frac{q_t}{\hat{q}_t} dx - \int \nabla q_t \cdot \nabla \log \frac{q_t}{\hat{q}_t} dx. \end{aligned}$$

Using $\nabla q_t = q_t \nabla \log q_t$ and splitting $\nabla \log q_t = \nabla \log(q_t/\hat{q}_t) + \nabla \log \hat{q}_t$:

$$- \int \nabla q_t \cdot \nabla \log \frac{q_t}{\hat{q}_t} = - \int q_t \left\| \nabla \log \frac{q_t}{\hat{q}_t} \right\|_2^2 dx - \int q_t \nabla \log \hat{q}_t \cdot \nabla \log \frac{q_t}{\hat{q}_t} dx.$$

Term (B). Substituting the FP equation for $\hat{q}_t$ and integrating by parts:

$$\begin{aligned} (B) &= \int \frac{q_t}{\hat{q}_t} \left(-\nabla \cdot (\hat{q}_t \hat{b}_t) + \Delta \hat{q}_t \right) dx \\ &= \int q_t \hat{b}_t \cdot \nabla \log \frac{q_t}{\hat{q}_t} dx - \int q_t \nabla \log \hat{q}_t \cdot \nabla \log \frac{q_t}{\hat{q}_t} dx, \end{aligned}$$

using $-\int \nabla \cdot (\hat{q}_t \hat{b}_t) \frac{q_t}{\hat{q}_t} = \int q_t \hat{b}_t \cdot \nabla \log \frac{q_t}{\hat{q}_t}$ and $\int \Delta \hat{q}_t \frac{q_t}{\hat{q}_t} = \int q_t \nabla \log \hat{q}_t \cdot \nabla \log \frac{q_t}{\hat{q}_t}$.

Combining. Substituting into (6.15), the $\nabla \log \hat{q}_t$ terms cancel, leaving

$$\frac{d}{dt} D_{\text{KL}}(q_t \parallel \hat{q}_t) = \int q_t \nabla \log \frac{q_t}{\hat{q}_t} \cdot (b_t - \hat{b}_t) dx - \int q_t \left\| \nabla \log \frac{q_t}{\hat{q}_t} \right\|_2^2 dx. \quad (6.16)$$

Applying Young's inequality $ab \leq \frac{\sigma^2}{2}a^2 + \frac{1}{2\sigma^2}b^2$ with $\sigma^2 = 2$ to the first integral:

$$\int q_t \nabla \log \frac{q_t}{\hat{q}_t} \cdot (b_t - \hat{b}_t) \leq \int q_t \left\| \nabla \log \frac{q_t}{\hat{q}_t} \right\|_2^2 dx + \frac{1}{4} \int q_t \|b_t - \hat{b}_t\|_2^2 dx.$$

The Fisher information terms cancel, giving $\frac{d}{dt} D_{\text{KL}}(q_t \| \hat{q}_t) \leq \frac{1}{4} \mathbb{E}_{q_t} \|b_t - \hat{b}_t\|_2^2$. Integrating from 0 to T gives (6.14). $\square$

Remark 6.6. Substituting the drifts (6.12), $b_t(y) - \hat{b}_t(y) = 2(\nabla \log p_{T-t}(y) - S_{T-t}(y, \theta))$, so the integrands in (6.13) and (6.14) are proportional to $\|\nabla \log p_{T-t}(y) - S_{T-t}(y, \theta)\|_2^2$, which corresponds exactly to the loss (6.10). Thus *minimizing the training loss directly controls the quality of the generated samples* in both W_2 and KL. For time discretizations, see for example [4, 23].

Week 3 Exercises

1. **(Probability flow ODE.)** The reverse SDE (6.4) has noise, which is useful for sampling but makes likelihood computation difficult. There is a deterministic counterpart, the *probability flow ODE*,

$$\frac{dY_t}{dt} = Y_t + \nabla \log p_{T-t}(Y_t), \quad Y_0 \sim p_T. \quad (6.17)$$

It is a purely deterministic flow. The key claim is that it shares the same marginal densities as the reverse SDE.

- (a) *(Same marginals as the reverse SDE.)* Let q_t^{ODE} denote the law of Y_t solving (6.17). Show that $q_t^{\text{ODE}} = p_{T-t}$ for all t by verifying that p_{T-t} satisfies the continuity equation for (6.17):

$$\partial_t p_{T-t} + \nabla \cdot (p_{T-t}(y + \nabla \log p_{T-t})) = 0.$$

Hint. Differentiate p_{T-t} in t using the Fokker–Planck equation for the forward OU process (6.2), and use the score trick $\Delta p = \nabla \cdot (p \nabla \log p)$.

- (b) *(Comparison with SDE.)* The reverse SDE (6.4) has drift $y + 2\nabla \log p_{T-t}(y)$, while the probability flow ODE (6.17) has drift $y + \nabla \log p_{T-t}(y)$. Both yield $q_T = p_0 = \pi$ when initialized at $q_0 = p_T$.
- (i) Explain why the ODE formulation allows exact likelihood computation via the instantaneous change-of-variables formula $\frac{d}{dt} \log q_t^{\text{ODE}}(Y_t) = -\nabla \cdot (Y_t + \nabla \log p_{T-t}(Y_t))$, while the SDE formulation does not.
 - (ii) Explain why the ODE formulation may allow faster generation.

2. **(Bounding the initialization error.)** The error bounds (6.13) and (6.14) each contain an initialization term measuring how well $\hat{q}_0 = \mathcal{N}(0, I_d)$ approximates $q_0 = p_T$.

- (a) *(W_2 initialization bound.)* The OU forward process (6.2) is Langevin dynamics (4.1) with $V(x) = \frac{1}{2}\|x\|^2$, so $\alpha = 1$ and stationary distribution $\mathcal{N}(0, I_d)$. Apply Theorem 4.6 directly to obtain

$$W_2(p_T, \mathcal{N}(0, I_d)) \leq e^{-T} W_2(\pi, \mathcal{N}(0, I_d)).$$

- (b) *(KL initialization bound.)* For the KL bound (6.14), we need $D_{\text{KL}}(q_0 \| \hat{q}_0) = D_{\text{KL}}(p_T \| \mathcal{N}(0, I_d))$. Since the OU process (6.2) is Langevin with stationary distribution $\gamma = \mathcal{N}(0, I_d)$, the KL decay identity from Week 2 Exercise 2(c)(i) gives

$$\frac{d}{dt} D_{\text{KL}}(p_t \| \gamma) = -\mathcal{I}(p_t \| \gamma), \quad \mathcal{I}(p_t \| \gamma) = \int \|\nabla \log(p_t / \gamma)\|_2^2 p_t dx.$$

(i) For $\gamma = \mathcal{N}(0, I_d)$, show that $\nabla \log \gamma(x) = -x$, so

$$\mathcal{I}(p_t \parallel \gamma) = \mathbb{E}_{p_t} \|\nabla \log p_t(X) + X\|_2^2.$$

(ii) We can show (accept without proof) that

$$D_{\text{KL}}(p_t \parallel \gamma) \leq \frac{1}{2} \mathcal{I}(p_t \parallel \gamma).$$

This is the *log-Sobolev inequality* for $\mathcal{N}(0, I_d)$. Combine with the decay identity to show $\frac{d}{dt} D_{\text{KL}}(p_t \parallel \gamma) \leq -2 D_{\text{KL}}(p_t \parallel \gamma)$, and conclude by Gronwall that

$$D_{\text{KL}}(p_T \parallel \mathcal{N}(0, I_d)) \leq e^{-2T} D_{\text{KL}}(\pi \parallel \mathcal{N}(0, I_d)).$$

(c) (*Complete W_2 error bound.*) Combine part (a) with (6.13). Let $\varepsilon_{\text{score}}^2 = \int_0^T \mathbb{E}_{q_t} [\|b_t(Y_t) - \hat{b}_t(Y_t)\|_2^2] dt$. Show that

$$W_2(\pi, \hat{q}_T) \leq e^{(\hat{L} + \frac{1}{2})T} (e^{-T} W_2(\pi, \mathcal{N}(0, I_d)) + \varepsilon_{\text{score}}).$$

What happens to this upper bound as $T \rightarrow \infty$?

(d) (*Complete KL error bound.*) Similarly combine part (b) with (6.14) to obtain

$$D_{\text{KL}}(\pi \parallel \hat{q}_T) \leq e^{-2T} D_{\text{KL}}(\pi \parallel \mathcal{N}(0, I_d)) + \int_0^T \frac{1}{4} \mathbb{E}_{q_t} [\|b_t(Y_t) - \hat{b}_t(Y_t)\|_2^2] dt.$$

Note that unlike the W_2 bound, this KL bound has *no exponentially growing prefactor* in T .

3. (Gaussian mixture: explicit computations.) Let $d = 1$ and $\pi = \frac{1}{2} \mathcal{N}(\mu, 1) + \frac{1}{2} \mathcal{N}(-\mu, 1)$ for $\mu > 0$.

(a) (*Marginals of the OU process.*) Show that

$$p_t = \frac{1}{2} \mathcal{N}(e^{-t} \mu, 1) + \frac{1}{2} \mathcal{N}(-e^{-t} \mu, 1).$$

(b) (*Explicit score.*) Write $w_t(x) = \phi(x; e^{-t} \mu, 1) / [\phi(x; e^{-t} \mu, 1) + \phi(x; -e^{-t} \mu, 1)]$ for the mixture weight, where $\phi(x; m, \sigma^2)$ is the $\mathcal{N}(m, \sigma^2)$ density. Show that

$$\nabla \log p_t(x) = -(x - e^{-t} \mu (2w_t(x) - 1)),$$

and verify that as $t \rightarrow \infty$, $w_t(x) \rightarrow \frac{1}{2}$ and $\nabla \log p_t(x) \rightarrow -x$ (score of $\mathcal{N}(0, 1)$).

7 Flow Matching, Rectified Flow, Stochastic Interpolants

7.1 Motivation: Beyond Gaussian Noise Source and Infinite Time

The score-based diffusion model of Section 6 has two structural limitations:

- ① *Finite-time truncation.* The forward OU process (6.2) reaches $\mathcal{N}(0, I_d)$ only as $T \rightarrow \infty$; in practice one truncates to $[0, T]$, introducing initialization errors.
- ② *Restricted transport.* The noising process always moves data toward Gaussian noise. One can therefore only learn to transport between a data distribution and $\mathcal{N}(0, I_d)$, not between two arbitrary distributions π_0 and π_1 .

In many applications, one wants a direct transport between two non-Gaussian distributions — for instance, π_0 a low-resolution image distribution and π_1 a high-resolution image distribution. Flow matching [14], as well as concurrent work, stochastic interpolants [2, 1] and rectified flows [16], addresses both issues by replacing the noising process with a general interpolation over a finite time interval.

7.2 Setup and Interpolation Process

Let $X_0 \sim \pi_0$ and $X_1 \sim \pi_1$ be samples from two target distributions (which may be accessed only through data), and let $Z \sim \mathcal{N}(0, I_d)$ be independent auxiliary noise. Define an *interpolation process*

$$I_t = I_t(X_0, X_1, Z), \quad t \in [0, 1], \quad (7.1)$$

satisfying the boundary conditions $I_0 = X_0$ and $I_1 = X_1$. Let μ_t denote the law of I_t (with density also written μ_t by abuse of notation), so $\mu_0 = \pi_0$ and $\mu_1 = \pi_1$. The goal is to construct a vector field b_t such that the ODE $dY_t = b_t(Y_t) dt$ transports π_0 to π_1 by matching the marginals of the interpolation.

Example 7.1 (Linear interpolant). *The simplest choice is the linear interpolant*

$$I_t = \alpha_t X_0 + \beta_t X_1, \quad \dot{I}_t = \dot{\alpha}_t X_0 + \dot{\beta}_t X_1, \quad (7.2)$$

where $\alpha_0 = \beta_1 = 1$ and $\alpha_1 = \beta_0 = 0$, so that $I_0 = X_0$ and $I_1 = X_1$. The canonical choice $\alpha_t = 1 - t$, $\beta_t = t$ gives $I_t = (1 - t)X_0 + tX_1$ with $\dot{I}_t = X_1 - X_0$.

Remark 7.2 (Connection to diffusion models). Consider an infinite time interval and $X_1 = Z \sim \mathcal{N}(0, I_d)$ and the interpolant coefficients $\alpha_t = e^{-t}$, $\beta_t = \sqrt{1 - e^{-2t}}$. By Theorem 6.2, the OU marginal satisfies $p_t(\cdot \mid x_0) = \mathcal{N}(e^{-t}x_0, (1 - e^{-2t})I_d)$, so

$$I_t \stackrel{d}{=} e^{-t}X_0 + \sqrt{1 - e^{-2t}} Z \sim p_t,$$

where $p_t = \text{Law}(X_t)$ is the OU forward process marginal. Thus $\mu_t = p_t$: the interpolant marginals coincide with the diffusion model forward marginals. See exercise at the end of this section.

Remark 7.3. In the above two examples, $I_t = I_t(X_0, X_1)$ without dependence on the latent random variable Z . This means that conditional on X_0, X_1 , I_t is a point mass. One can introduce additional stochasticity so that $I_t|X_0, X_1$ has some variability and there is a conditional density $p_t(\cdot|X_0, X_1)$. Using the latent Z can achieve this. In the original paper on flow matching, the conditional distribution is directly modeled, while in stochastic interpolants, the interpolant process is the primary modeling object.

7.3 The Matching Theorem

Theorem 7.4. *Let I_t be an interpolation process (7.1) with marginal law μ_t , and define the conditional velocity field*

$$b_t(x) = \mathbb{E}[\dot{I}_t | I_t = x], \quad \dot{I}_t := \frac{d}{dt}I_t. \quad (7.3)$$

Let Y_t solve the ODE

$$dY_t = b_t(Y_t) dt, \quad Y_0 \sim \pi_0. \quad (7.4)$$

Then $\text{Law}(Y_t) = \mu_t$ for all $t \in [0, 1]$. In particular, $Y_1 \sim \pi_1$.

Proof. Assume μ_t has a smooth density (also denoted μ_t); otherwise one uses a weak formulation. For any test function $\phi \in C_c^\infty(\mathbb{R}^d)$, differentiate $\int \mu_t \phi dx = \mathbb{E}[\phi(I_t)]$:

$$\int \partial_t \mu_t \phi dx = \frac{d}{dt} \mathbb{E}[\phi(I_t)] = \mathbb{E}[\nabla \phi(I_t) \cdot \dot{I}_t].$$

By the tower property of conditional expectation,

$$\begin{aligned} \mathbb{E}[\nabla \phi(I_t) \cdot \dot{I}_t] &= \mathbb{E}[\nabla \phi(I_t) \cdot \mathbb{E}[\dot{I}_t | I_t]] = \int \nabla \phi(x) \cdot b_t(x) \mu_t(x) dx \\ &= - \int \phi(x) \nabla \cdot (\mu_t b_t)(x) dx, \end{aligned}$$

where the last equality is integration by parts. Since ϕ is arbitrary:

$$\partial_t \mu_t + \nabla \cdot (\mu_t b_t) = 0. \quad (7.5)$$

This is the continuity equation (3.3) for the ODE (7.4). Since $\text{Law}(Y_0) = \pi_0 = \mu_0$ and both $\text{Law}(Y_t)$ and μ_t satisfy (7.5) with the same initial condition, they coincide for all t . $\square$

7.4 Algorithm: Training by Matching the Velocity

Theorem 7.4 reduces transporting π_0 to π_1 to learning $b_t(x) = \mathbb{E}[\dot{I}_t | I_t = x]$. Since the conditional expectation is the L^2 -projection, b_t is the minimizer of the *flow matching loss*:

$$L(\theta) = \int_0^1 \mathbb{E}[\|\hat{b}_t(I_t, \theta) - \dot{I}_t\|_2^2] dt. \quad (7.6)$$

The loss (7.6) is fully tractable: at each step, sample $X_0 \sim \pi_0$, $X_1 \sim \pi_1$, and $Z \sim \mathcal{N}(0, I_d)$ from data; evaluate I_t and $\dot{I}_t$ in closed form from (7.1); and optimize with SGD/ADAM. Once trained, generate new samples from π_1 by integrating the ODE (7.4) from $Y_0 \sim \pi_0$.

Error bound. Applying the same coupling argument as Theorem 6.4, if $\hat{b}_t$ is $\hat{L}$ -Lipschitz in x uniformly in t , then

$$W_2^2(\hat{\mu}_1, \mu_1) \leq e^{(2\hat{L}+1)} \int_0^1 \mathbb{E}[\|\hat{b}_t(I_t, \theta) - b_t(I_t)\|_2^2] dt, \quad (7.7)$$

where $\hat{\mu}_1$ is the law of $\hat{Y}_1$ from the approximate ODE $d\hat{Y}_t = \hat{b}_t(\hat{Y}_t, \theta) dt$. Note the absence of the initialization error term $W_2(q_0, \hat{q}_0)^2$: with exact boundary conditions $I_0 = Z \sim \mathcal{N}(0, I_d) = \hat{Y}_0$, the initialization error is zero.

Remark 7.5 (Summary of advantages over diffusion models). For the same interpolant choice as in Theorem 7.2, flow matching and diffusion models share the same forward marginals $\mu_t = p_t$. Flow matching's advantages are:

- *Finite horizon:* the time interval is $[0, 1]$, with exact boundary conditions and no truncation error.
- *General transport:* π_0 and π_1 can be any two distributions, not just data and Gaussian noise.
- *Simpler framework:* the loss (7.6) regresses onto $\dot{I}_t$, which is known in closed form. In comparison, denoising score matching (6.9) regresses onto the conditional score $\nabla_{X_t} \log p_t(X_t | X_0)$ whose form requires some calculations to obtain.

7.5 Connection to Score Based Diffusion Models

In this subsection we specialize to the interpolant $I_t = \alpha_t Z + \beta_t X_1$ which goes from Z at time $t = 0$ to X_1 at $t = 1$, and derive explicit formulas connecting b_t , the score $\nabla \log \mu_t$, and the diffusion model.

Noise predictor and denoiser. Taking conditional expectations in $I_t = \alpha_t Z + \beta_t X_1$ gives:

$$x = \underbrace{\alpha_t \mathbb{E}[Z | I_t = x]}_{\text{noise predictor}} + \underbrace{\beta_t \mathbb{E}[X_1 | I_t = x]}_{\text{denoiser}}, \quad b_t(x) = \dot{\alpha}_t \mathbb{E}[Z | I_t = x] + \dot{\beta}_t \mathbb{E}[X_1 | I_t = x]. \quad (7.8)$$

Thus b_t can equivalently be expressed through the noise predictor $\mathbb{E}[Z | I_t = x]$ or the denoiser $\mathbb{E}[X_1 | I_t = x]$, analogously to (6.8) in the diffusion model.

Score of the interpolant marginals. The following lemma is the key tool.

Lemma 7.6 (Score via conditional expectation). *Let X, Y be jointly distributed random variables with p the density of X . Then*

$$\nabla \log p(x) = \mathbb{E}[\nabla_X \log p(X | Y) | X = x]. \quad (7.9)$$

Proof. Differentiating $p(x) = \int p(x | y) q(y) dy$: $\nabla_x p(x) = \int \nabla_x p(x | y) q(y) dy$. Dividing by $p(x)$ and recognizing the posterior:

$$\nabla \log p(x) = \int \nabla_x \log p(x | y) \frac{p(x | y) q(y)}{p(x)} dy = \mathbb{E}[\nabla_X \log p(X | Y) | X = x].$$

□

Proposition 7.7 (Score and noise predictor). *For the interpolant $I_t = \alpha_t Z + \beta_t X_1$,*

$$\nabla \log \mu_t(x) = -\frac{1}{\alpha_t} \mathbb{E}[Z | I_t = x]. \quad (7.10)$$

Proof. Apply Theorem 7.6 with $X = I_t$, $Y = X_1$. Given $X_1 = x_1$, $I_t | X_1 = x_1 \sim \mathcal{N}(\beta_t x_1, \alpha_t^2 I_d)$, so

$$\nabla_x \log p_{t|1}(x | x_1) = -\frac{x - \beta_t x_1}{\alpha_t^2}.$$

By Theorem 7.6:

$$\nabla \log \mu_t(x) = \mathbb{E}\left[-\frac{I_t - \beta_t X_1}{\alpha_t^2} \mid I_t = x\right] = -\frac{1}{\alpha_t^2} \mathbb{E}[I_t - \beta_t X_1 | I_t = x] = -\frac{1}{\alpha_t^2} \mathbb{E}[\alpha_t Z | I_t = x].$$

□

Solving (7.8) and (7.10) for $\mathbb{E}[Z | I_t = x]$ and substituting into the expression for b_t :

$$b_t(x) = \alpha_t^2 \left(\frac{\dot{\beta}_t}{\beta_t} - \frac{\dot{\alpha}_t}{\alpha_t} \right) \nabla \log \mu_t(x) + \frac{\dot{\beta}_t}{\beta_t} x. \quad (7.11)$$

This expresses the flow matching velocity as an affine function of the score $\nabla \log \mu_t$, mirroring (4.3) in the Langevin setting.

A family of SDEs with the same marginals. By the score identity $\nabla \cdot (\mu_t \nabla \log \mu_t) = \Delta \mu_t$, adding any diffusion $g_t \geq 0$ to the drift of (7.4) preserves the marginals.

Proposition 7.8. *For any $g_t \geq 0$, the SDE*

$$dY_t = (b_t(Y_t) + g_t \nabla \log \mu_t(Y_t)) dt + \sqrt{2g_t} dW_t, \quad Y_0 \sim \mathcal{N}(0, I_d), \quad (7.12)$$

satisfies $\text{Law}(Y_t) = \mu_t$ for all $t \in [0, 1]$.

Proof. The Fokker–Planck equation for (7.12) is

$$\partial_t \mu_t + \nabla \cdot (\mu_t (b_t + g_t \nabla \log \mu_t)) = g_t \Delta \mu_t.$$

Since $\nabla \cdot (\mu_t \nabla \log \mu_t) = \Delta \mu_t$, the right-hand side equals $g_t \nabla \cdot (\mu_t \nabla \log \mu_t)$, which cancels the added term, reducing to the continuity equation (7.5) that μ_t satisfies. □

The ODE (7.4) corresponds to $g_t = 0$. Equation (7.11) shows that (7.12) is always an SDE whose drift is an affine function of the score $\nabla \log \mu_t$.

Remark 7.9 (Stochastic interpolants: score formula). For the stochastic interpolant $I_t = \alpha_t X_0 + \beta_t X_1 + \sigma_t Z$ with $\sigma_0 = \sigma_1 = 0$, Theorem 7.6 applied with $Y = (X_0, X_1)$ gives $\nabla \log \mu_t(x) = -\mathbb{E}[Z/\sigma_t | I_t = x]$, and $b_t(x) = \mathbb{E}[\dot{\alpha}_t X_0 + \dot{\beta}_t X_1 + \dot{\sigma}_t Z | I_t = x]$. The family (7.12) applies unchanged.

7.6 Rectified Flow and Optimal Transport

Straight trajectories and numerical efficiency. For the linear interpolant $I_t = \alpha_t X_0 + \beta_t X_1$ with the rectified flow choice $\alpha_t = 1 - t$, $\beta_t = t$, the conditional velocity is $b_t(x) = \mathbb{E}[X_1 - X_0 \mid I_t = x]$. If trajectories of the ODE (7.4) were perfectly straight — i.e., $b_t(Y_t) = Y_1 - Y_0$ along each trajectory — then the ODE could be solved exactly with a single Euler step. Trajectory crossing prevents this in general: when (X_0, X_1) are drawn independently from $\pi_0 \times \pi_1$, different pairs (X_0, X_1) and (X'_0, X'_1) can produce paths I_t and I'_t that intersect at some time t , forcing b_t at the crossing point to average over multiple directions.

The reflow procedure. Given a trained velocity field $\hat{b}_t$, one can generate a new *straightened* coupling by simulating the learned ODE: starting from $X_0 \sim \pi_0$, run

$$\frac{dY_t}{dt} = \hat{b}_t(Y_t), \quad Y_0 = X_0,$$

and set $X'_1 := Y_1$. By construction, (X_0, X'_1) is a new coupling of π_0 and π_1 (since $Y_1 \sim \pi_1$ when the ODE is exact), and the straight-line paths $(1 - t)X_0 + tX'_1$ are precisely the ODE trajectories. These cannot cross, so training a new flow on this coupling produces a velocity field that is closer to a pure translation $b_t(x) \approx X'_1 - X_0$ at each point. Iterating this *reflow* [16] procedure

$$(X_0^{(k+1)}, X_1^{(k+1)}) := \text{ODE coupling from } \hat{b}_t^{(k)}$$

progressively straightens the trajectories.

Connection to optimal transport. The *Monge optimal transport* (OT) problem between π_0 and π_1 seeks a transport map $T : \mathbb{R}^d \rightarrow \mathbb{R}^d$ with $T_{\#}\pi_0 = \pi_1$ minimizing the quadratic cost:

$$\min_{T: T_{\#}\pi_0 = \pi_1} \mathbb{E}_{X_0 \sim \pi_0} [\|X_0 - T(X_0)\|_2^2]. \quad (7.13)$$

Under the OT coupling (X_0, X_1^*) where $X_1^* = T(X_0)$, the straight-line paths $(1 - t)X_0 + tT(X_0)$ never cross (a consequence of cyclical monotonicity of the OT map), so the velocity $b_t(x) = T(X_0) - X_0$ is a pure translation and the flow ODE can be solved in one step.

The reflow procedure approximates this: each iteration reduces trajectory crossings, and in the limit the coupled pairs $(X_0^{(\infty)}, X_1^{(\infty)})$ converge to the OT coupling. The rectified flow therefore provides a practical algorithm in high dimensions for approximating the OT map without solving (7.13) directly. For related diffusion methods for approximating the entropy regularized OT, namely Schrödinger’s bridge, see [20].

Week 4 Exercises

1. (**W_2 error bound for flow matching.**) Prove the error bound (7.7) by adapting the coupling argument of Theorem 6.4.

(a) (*Setup.*) Let Y_t solve the true ODE $dY_t = b_t(Y_t) dt$, $Y_0 \sim \pi_0$, and $\hat{Y}_t$ solve the approximate ODE $d\hat{Y}_t = \hat{b}_t(\hat{Y}_t, \theta) dt$, $\hat{Y}_0 \sim \pi_0$. Since both are initialized at the same π_0 , one can take $Y_0 = \hat{Y}_0$ (i.e. couple them at initialization). Compute $\frac{d}{dt} \|Y_t - \hat{Y}_t\|_2^2$ and apply the same add-and-subtract and Young/Lipschitz argument as in Theorem 6.4 to obtain

$$\frac{d}{dt} \|Y_t - \hat{Y}_t\|_2^2 \leq (2\hat{L} + 1) \|Y_t - \hat{Y}_t\|_2^2 + \|b_t(Y_t) - \hat{b}_t(Y_t)\|_2^2.$$

(b) (*Gronwall and initialization.*) Apply Gronwall to the inequality in (a) and take expectations to show

$$\mathbb{E} \|Y_1 - \hat{Y}_1\|_2^2 \leq e^{(2\hat{L}+1)} \int_0^1 \mathbb{E} [\|b_t(Y_t) - \hat{b}_t(Y_t)\|_2^2] dt.$$

Explain why there is no initialization error term $W_2(\mu_0, \hat{\mu}_0)^2$ here, in contrast to (6.13).

The diffusion model bound (6.13) has a prefactor $e^{(2\hat{L}+1)T}$ that grows with the time horizon $T \rightarrow \infty$. The flow matching bound (7.7) has prefactor $e^{(2\hat{L}+1)}$ (with $T = 1$ fixed). Nevertheless, the Lipschitz constant $\hat{L}$ can be large in practice.

2. (**Infinite-time interpolant and score-based diffusion models.**) This exercise establishes the precise relationship between the flow matching framework and the diffusion model of Section 6.

In Section 7 the interpolant is $I_t = \alpha_t Z + \beta_t X_1$ with $Z \sim \mathcal{N}(0, I_d)$ (source/noise) and $X_1 \sim \pi_1$ (target). For the diffusion model, the source is data $X_0 \sim \pi_0$ and the target is Gaussian noise $Z' \sim \mathcal{N}(0, I_d)$ — the roles are swapped. So set $\pi_0 = \pi$ (data), $\pi_1 = \mathcal{N}(0, I_d)$, and consider the interpolant with the data as source:

$$I_t = \alpha_t X_0 + \beta_t Z', \quad X_0 \sim \pi_0, \quad Z' \sim \mathcal{N}(0, I_d) \text{ independent}, \quad (7.14)$$

with $\alpha_t = e^{-t}$, $\beta_t = \sqrt{1 - e^{-2t}}$, $\alpha_0 = 1$, $\beta_0 = 0$, $\alpha_\infty = 0$, $\beta_\infty = 1$. This matches Theorem 6.2: $I_t \stackrel{d}{=} X_t$ where X_t solves the forward OU (6.2), so $\mu_t = p_t$.

(a) (*Score in the swapped convention.*) In the convention of (7.14), X_0 plays the role of Z and Z' plays the role of X_1 . Apply Theorem 7.6 with $X = I_t$, $Y = X_0$: given $X_0 = x_0$, $I_t | X_0 \sim \mathcal{N}(\alpha_t x_0, \beta_t^2 I_d)$, so $\nabla_x \log p_{t|0}(x | x_0) = -(x - \alpha_t x_0) / \beta_t^2$. Show that

$$\nabla \log \mu_t(x) = \nabla \log p_t(x) = -\frac{1}{\beta_t} \mathbb{E}[Z' | I_t = x]. \quad (7.15)$$

This is the analogue of (7.10) with $Z \leftrightarrow X_0$ and $X_1 \leftrightarrow Z'$: the “noise predictor” is now $\mathbb{E}[Z' \mid I_t = x]$.

- (b) (*Velocity field b_t and the OU drift.*) The conditional velocity is $b_t(x) = \dot{\alpha}_t \mathbb{E}[X_0 \mid I_t = x] + \dot{\beta}_t \mathbb{E}[Z' \mid I_t = x]$. Using (7.15) and $x = \alpha_t \mathbb{E}[X_0 \mid I_t = x] + \beta_t \mathbb{E}[Z' \mid I_t = x]$, eliminate $\mathbb{E}[X_0 \mid I_t = x]$ and $\mathbb{E}[Z' \mid I_t = x]$ to show

$$b_t(x) = \frac{\dot{\alpha}_t}{\alpha_t} x + \beta_t \left(\frac{\dot{\alpha}_t \beta_t}{\alpha_t} - \dot{\beta}_t \right) \nabla \log p_t(x). \quad (7.16)$$

Substituting $\alpha_t = e^{-t}$, $\beta_t = \sqrt{1 - e^{-2t}}$, show that $\dot{\alpha}_t/\alpha_t = -1$ and $\dot{\alpha}_t \beta_t/\alpha_t - \dot{\beta}_t = -\beta_t - e^{-2t}/\beta_t = -(1 - e^{-2t} + e^{-2t})/\beta_t = -1/\beta_t$, giving

$$b_t(x) = -x - \nabla \log p_t(x).$$

- (c) (*Recovering the forward OU and reverse SDE.*) Using Theorem 7.8 with $b_t(x) = -x - \nabla \log p_t(x)$:

- (i) Show that $g_t = 1$ gives drift $b_t + \nabla \log \mu_t = -x$ and SDE $dY_t = -Y_t dt + \sqrt{2} dW_t$, which is exactly the forward OU process (6.2). This confirms the family (7.12) contains the forward OU as the $g_t = 1$ member.
- (ii) Apply Anderson’s time-reversal (Theorem 6.3) to the forward OU process to obtain the reverse SDE (6.4). Since any member of (7.12) shares the marginals p_t , their time-reversals all yield the same marginal law p_{T-t} for $Y_t = X_{T-t}$. In particular, note that the flow matching ODE (7.4) with $b_t(x) = -x - \nabla \log p_t(x)$ has the same marginals p_t as the forward OU SDE. Its time-reversal $\tilde{Y}_s = Y_{T-s}$ satisfies $d\tilde{Y}_s = -b_{T-s}(\tilde{Y}_s) ds = (X_{T-s} + \nabla \log p_{T-s}(\tilde{Y}_s)) ds$, i.e. $d\tilde{Y}_s = (\tilde{Y}_s + \nabla \log p_{T-s}(\tilde{Y}_s)) ds$. This matches the probability flow ODE (6.17).

As a summary, the forward OU SDE, the flow matching ODE, and the SDE family (7.12) all share marginals p_t ; Anderson’s time-reversal of any of them with diffusion $\sqrt{2}$ gives the reverse SDE (6.4), while reversing the ODE ($g_t = 0$) gives the probability flow ODE (6.17).

8 Conditional Sampling with Generative Priors

8.1 Motivation and Setting

The preceding sections developed generative models that learn to transport noise to a data distribution π . Given unlabeled data $x^{(1)}, \dots, x^{(N)} \sim \pi$, we trained SDEs/ODEs of the form

$$dX_t = b_t(X_t) dt + \sigma_t dW_t, \quad X_0 \sim \mathcal{N}(0, I_d), \quad X_1 \sim \pi,$$

where for example (i) $b_t(x) = x + 2\nabla \log p_t(x)$, $\sigma_t = \sqrt{2}$ (reverse SDE), or (ii) $b_t(x) = \mathbb{E}[\dot{I}_t | I_t = x]$, $\sigma_t = 0$ (flow matching ODE). In both cases the generative model samples *unconditionally* from π . Note that in case (ii), we can make it SDE by adding score terms in the drift, see Proposition 7.8.

In practice, one often wants to *control* the generated samples. Two canonical settings arise:

- ① **Bayesian inverse problem.** One has unlabeled data $x^{(1)}, \dots, x^{(N)} \sim \pi$ and a *forward model* $y = F(x) + \sigma\varepsilon$, $\varepsilon \sim \mathcal{N}(0, I)$, producing a noisy observation $y \in \mathbb{R}^m$. The goal is to sample the *posterior*

$$P(x | y) \propto \pi(x) \exp\left(-\frac{\|y - F(x)\|_2^2}{2\sigma^2}\right).$$

The likelihood $P(y | x) \propto \exp(-\|y - F(x)\|^2/(2\sigma^2))$ is known but $\pi(x)$ is known only through samples.

- ② **Reward tilting / fine-tuning / inference-time scaling.** One wants to sample from $\pi(x) \exp(R(x; y))$ for a *reward function* $R : \mathbb{R}^d \times \mathbb{R}^m \rightarrow \mathbb{R}$, where y is a conditioning variable. Intuition: generate samples that sit in the region where the reward R is large while staying close to the prior π . We can define a likelihood $P(y|x) \propto \exp(R(x; y))$.

Both settings reduce to the same problem: given a generative model for π and a likelihood $P(y | x)$, sample from $P(x | y) \propto \pi(x) P(y | x)$ for a fixed observation y . We present two alternative approaches with different features: *conditional flow matching* (Setting 1, requires paired training data and retraining) and *guidance* (Setting 2, uses a pre-trained unconditional model with a tilting term at inference time).

8.2 Setting 1: Amortized Conditional Flow Matching

Key idea. Suppose we can generate *paired* samples by first drawing $X_1 \sim \pi$ and then drawing $Y | X_1 \sim P(y | X_1)$. We use these pairs to train a new conditional velocity field $b_t(x, y)$ that transports noise to $P(\cdot | y)$.

Setup. Let $X_0 \sim \mathcal{N}(0, I_d)$ and $X_1 \sim \pi$ be independent, and draw $Y \mid X_1 \sim P(y \mid X_1)$. Consider the linear interpolant $I_t = \alpha_t X_0 + \beta_t X_1$ with $I_0 = X_0$, $I_1 = X_1$. Define the *conditional velocity field*

$$b_t(x, y) := \mathbb{E}[\dot{I}_t \mid I_t = x, Y = y], \quad \dot{I}_t = \dot{\alpha}_t X_0 + \dot{\beta}_t X_1. \quad (8.1)$$

Let $p_t(x \mid y)$ denote the conditional density of I_t given $Y = y$, so $p_0(\cdot \mid y) = \mathcal{N}(0, I_d)$ for all y (since $I_0 = X_0 \perp Y$) and $p_1(\cdot \mid y) \propto \pi(x) P(y \mid x)$.

Proposition 8.1 (Conditional flow matching). *Let $\tilde{X}_t$ solve the ODE*

$$d\tilde{X}_t = b_t(\tilde{X}_t, y) dt, \quad \tilde{X}_0 \sim \mathcal{N}(0, I_d), \quad (8.2)$$

for a fixed y . Then $\text{Law}(\tilde{X}_t) = p_t(\cdot \mid y)$ for all $t \in [0, 1]$. In particular, $\tilde{X}_1 \sim P(\cdot \mid y) \propto \pi(x) P(y \mid x)$.

Proof. Fix y and let $\phi \in C_c^\infty(\mathbb{R}^d)$. Differentiating $\int p_t(x \mid y) \phi(x) dx = \mathbb{E}[\phi(I_t) \mid Y = y]$ in t :

$$\int \partial_t p_t(x \mid y) \phi(x) dx = \frac{d}{dt} \mathbb{E}[\phi(I_t) \mid Y = y] = \mathbb{E}[\nabla \phi(I_t) \cdot \dot{I}_t \mid Y = y].$$

By the tower property, conditioning on (I_t, Y) inside the expectation conditioned on $Y = y$:

$$\begin{aligned} \mathbb{E}[\nabla \phi(I_t) \cdot \dot{I}_t \mid Y = y] &= \mathbb{E}[\nabla \phi(I_t) \cdot \mathbb{E}[\dot{I}_t \mid I_t, Y] \mid Y = y] \\ &= \int \nabla_x \phi(x) \cdot b_t(x, y) p_t(x \mid y) dx \\ &= - \int \phi(x) \nabla_x \cdot (p_t(x \mid y) b_t(x, y)) dx. \end{aligned}$$

Since ϕ is arbitrary, $p_t(\cdot \mid y)$ satisfies

$$\partial_t p_t(x \mid y) + \nabla_x \cdot (p_t(x \mid y) b_t(x, y)) = 0. \quad (8.3)$$

This is the continuity equation for (8.2) with the same initial condition $p_0(\cdot \mid y) = \mathcal{N}(0, I_d)$, so $\text{Law}(\tilde{X}_t) = p_t(\cdot \mid y)$ for all t . $\square$

Training. The conditional velocity $b_t(x, y)$ minimizes the loss

$$L(\theta) = \int_0^1 \mathbb{E}[\|\hat{b}_t(I_t, Y, \theta) - \dot{I}_t\|_2^2] dt, \quad (8.4)$$

where the expectation is over $X_0 \sim \mathcal{N}(0, I_d)$, $X_1 \sim \pi$, $Y \mid X_1 \sim P(\cdot \mid X_1)$, and $I_t = \alpha_t X_0 + \beta_t X_1$. Training proceeds by first drawing $I_t \mid X_0, X_1$, then $Y \mid X_1$, and regressing $\hat{b}_t(I_t, Y, \theta)$ onto $\dot{I}_t$.

Remark 8.2 (Amortized inference). The trained network $\hat{b}_t(x, y, \theta)$ can be evaluated at any y to generate approximate samples from $P(\cdot \mid y)$ by integrating (8.2). However, the loss (8.4) only guarantees accuracy of $P(x \mid y)$ averaged over $Y \sim P(y)$. For a specific fixed observation y , errors may be larger. This is *amortized inference*: one trained model serves all y simultaneously, at the cost of per-instance accuracy.

Remark 8.3 (Classifier-free guidance). In modern generative modeling practice, conditional flow matching is often combined with the unconditional model into a single network $\hat{b}_t(x, y, \theta)$ trained jointly on both conditional and unconditional samples (the latter obtained by replacing Y with a null token $\emptyset$). At inference time, one uses the extrapolated drift

$$(1 + w) \hat{b}_t(x, y, \theta) - w \hat{b}_t(x, \emptyset, \theta), \quad w \geq 0,$$

which heuristically amplifies the conditional signal. This is *classifier-free guidance* (Ho–Salimans 2022 [11]) and is a variant of Setting 1 that empirically gives sharper conditional samples for a fixed y at the cost of a bias controlled by the guidance strength w .

In the above setting, the information of $P(y|x)$ is not used directly; it is used based on the simulated many samples Y . This is also referred to as *simulation-based inference*.

8.3 Setting 2: Guidance without Re-training

Setup. Setting 2 takes a pre-trained unconditional generative model for π and steers it toward $P(\cdot | y)$ at inference time for a specific fixed observation y , without any retraining.

We work in the flow matching setting of Section 7. By (7.11), the unconditional velocity field satisfies

$$b_t(x) = \frac{\dot{\beta}_t}{\beta_t} x + \alpha_t^2 \left(\frac{\dot{\beta}_t}{\beta_t} - \frac{\dot{\alpha}_t}{\alpha_t} \right) \nabla_x \log p_t(x), \quad (8.5)$$

i.e. b_t depends on x through the score $\nabla_x \log p_t(x)$ and the linear term $\frac{\dot{\beta}_t}{\beta_t} x$. Combining (8.5) with the SDE family (7.12), the unconditional generative SDE has drift

$$b_t(x) + g_t \nabla_x \log p_t(x) = \underbrace{\frac{\dot{\beta}_t}{\beta_t} x + \left(\alpha_t^2 \left(\frac{\dot{\beta}_t}{\beta_t} - \frac{\dot{\alpha}_t}{\alpha_t} \right) + g_t \right) \nabla_x \log p_t(x)}_{=: c_t}. \quad (8.6)$$

The drift is an affine function of the score with coefficient c_t .

From unconditional to conditional via Bayes. To sample from $p_t(\cdot | y)$ instead of $p_t(\cdot)$, we replace $\nabla_x \log p_t(x)$ in (8.6) by $\nabla_x \log p_t(x | y)$. Bayes’ rule gives

$$\nabla_x \log p_t(x | y) = \underbrace{\nabla_x \log p_t(x)}_{\text{prior score (known)}} + \underbrace{\nabla_x \log p_t(y | x)}_{\text{likelihood score}}. \quad (8.7)$$

The prior score is available from the trained model. The likelihood score is the difficulty.

The marginal likelihood. The likelihood of observation y given $I_t = x$ is

$$p_t(y | x) = \int p_1(y | x_1) p_t(x_1 | x) dx_1, \quad (8.8)$$

where $p_1(y | x_1) = P(y | x_1)$ is the data likelihood and $p_t(x_1 | x)$ is the denoising distribution at level t . Computing (8.8) is generally intractable.

Guidance approximation. Approximate $p_t(\cdot | x)$ by a point mass at its conditional mean:

$$p_t(y | x) \approx p_1(y | \mathbb{E}[X_1 | I_t = x]). \quad (8.9)$$

Here $\mathbb{E}[X_1 | I_t = x]$ is the *denoiser* of the prior, which is available from the trained model via Theorem 7.7 (in particular, expressible through the prior score $\nabla_x \log p_t(x)$). Define the *guidance function*

$$f_t(x, y) := \nabla_x \log p_1(y | \mathbb{E}[X_1 | I_t = x]). \quad (8.10)$$

Then (8.7) is approximated by $\nabla_x \log p_t(x | y) \approx \nabla_x \log p_t(x) + f_t(x, y)$.

Guided generation. Substituting this approximation into (8.6) gives the *guided drift*:

$$\frac{\dot{\beta}_t}{\beta_t} x + c_t(\nabla_x \log p_t(x) + f_t(x, y)). \quad (8.11)$$

The guided process is

$$d\tilde{X}_t = \left[\frac{\dot{\beta}_t}{\beta_t} \tilde{X}_t + c_t(\nabla_x \log p_t(\tilde{X}_t) + f_t(\tilde{X}_t, y)) \right] dt + \sqrt{2g_t} dW_t, \quad \tilde{X}_0 \sim \mathcal{N}(0, I_d). \quad (8.12)$$

When the approximation (8.9) is exact, Theorem 7.8 (applied to the conditional marginals $p_t(\cdot | y)$) guarantees $\text{Law}(\tilde{X}_t) = p_t(\cdot | y)$ and $\tilde{X}_1 \sim P(\cdot | y)$.

Remark 8.4 (Diffusion model guidance). For the score-based diffusion model of Section 6 (the reverse SDE (6.4) with explicit drift $y + 2\nabla_x \log p_{T-t}(y)$), the guided version simply replaces the score by the guided score:

$$d\tilde{Y}_t = (\tilde{Y}_t + 2(\nabla_x \log p_{T-t}(\tilde{Y}_t) + f_{T-t}(\tilde{Y}_t, y))) dt + \sqrt{2} dW_t.$$

This is like *classifier guidance* (Dhariwal–Nichol 2021 [9]): the guidance term $f_t(x, y)$ approximates $\nabla_x \log p_t(y | x)$ via (8.9) (DPS [8]) or via a separately trained classifier $\hat{p}_t(y | x)$ [9].

Remark 8.5 (Bias of the guidance approximation). The approximation (8.9) introduces a bias: it is exact only when the denoising distribution $p_t(\cdot | x)$ is a point mass, which holds at $t = 1$ but not at intermediate t . The bias is most severe at small t , when $p_t(\cdot | x)$ is spread over multiple modes and the conditional mean $\mathbb{E}[X_1 | I_t = x]$ is a poor summary.

8.4 Discussions and Summary

Several lines of work aim to reduce or remove this bias:

- *Sequential Monte Carlo (SMC)*, *annealed importance sampling (AIS)* and *twisted particle filtering*: run multiple guided trajectories (particles) carrying importance weights, with periodic resampling to focus computation on high-weight particles.

The “twist” refers to the choice of guided drift that makes the particle proposals closer to the target conditional distribution; (8.10) is one such twist. With proper weighting, SMC yields asymptotically exact samples from $P(\cdot | y)$ as the number of particles grows. *This is the dynamical analogue of importance sampling from Section 2:* the guided SDE plays the role of the proposal g , and the importance weights play the role of π/g , but now along an entire path rather than at a single point. The dynamical structure is what makes the proposal close to the target, mitigating the curse of dimensionality that crippled static IS (recall Theorem 2.10).

- *Stochastic optimal control:* formulate conditional sampling as an optimal control problem whose solution is the exact conditional drift, with the guidance term (8.10) arising as a first-order approximation to the optimal control.

These methods trade additional computation at inference time for reduced bias in the conditional samples.

Remark 8.6 (Comparison of the two settings). We summarize as follows.

- *Setting 1 (conditional flow matching):* requires paired training data (X_1, Y) and retraining the model; generalizes over all y in an amortized way; classifier-free guidance is a popular variant that gives sharper conditional samples.
- *Setting 2 (guidance):* uses only a pre-trained unconditional model and a differentiable likelihood $P(y | x)$; requires no retraining; works for a specific y at inference time; the basic form has bias controlled by (8.9), which can be reduced by SMC or control-theoretic corrections.

Neither setting strictly dominates the other: Setting 1 invests upfront training cost for fast conditional sampling; Setting 2 invests inference-time cost for flexibility and biased but cheap conditional sampling.

Week 5 Exercises

1. **(Error bound for conditional flow matching.)** This exercise derives an error bound analogous to (7.7) for the conditional flow matching of Section 8.2, making precise the amortization claim that the trained model only guarantees accuracy of $P(\cdot | y)$ averaged over Y .

Setup. Recall the conditional velocity field $b_t(x, y) = \mathbb{E}[\dot{I}_t | I_t = x, Y = y]$ from (8.1). Let $\hat{b}_t(x, y, \theta)$ be a trained network (not necessarily the minimizer of (8.4)). For each fixed y , let $\tilde{X}_t^y$ solve the true ODE $d\tilde{X}_t^y = b_t(\tilde{X}_t^y, y) dt$ with $\tilde{X}_0^y \sim \mathcal{N}(0, I_d)$, and let $\hat{X}_t^y$ solve the approximate ODE $d\hat{X}_t^y = \hat{b}_t(\hat{X}_t^y, y, \theta) dt$ with $\hat{X}_0^y \sim \mathcal{N}(0, I_d)$. By Theorem 8.1, $\tilde{X}_1^y \sim P(\cdot | y)$. Denote $\hat{P}(\cdot | y) = \text{Law}(\hat{X}_1^y)$.

- (a) *(Pointwise bound for fixed y .)* For a fixed observation y , assume $\hat{b}_t(\cdot, y, \theta)$ is $\hat{L}$ -Lipschitz in x uniformly in t . Adapt the coupling argument of Theorem 6.4 (with $Y_0 = \hat{Y}_0$ exact, since both ODEs start from $\mathcal{N}(0, I_d)$) to show

$$W_2^2(\hat{P}(\cdot | y), P(\cdot | y)) \leq e^{(2\hat{L}+1)} \int_0^1 \mathbb{E} \left[\|\hat{b}_t(I_t, y, \theta) - b_t(I_t, y)\|_2^2 \mid Y = y \right] dt.$$

- (b) *(Bound averaged over Y .)* Take expectation over $Y \sim P(y) = \int \pi(x_1) P(y | x_1) dx_1$ in part (a) to obtain

$$\mathbb{E}_Y [W_2^2(\hat{P}(\cdot | Y), P(\cdot | Y))] \leq e^{(2\hat{L}+1)} \int_0^1 \mathbb{E} \left[\|\hat{b}_t(I_t, Y, \theta) - b_t(I_t, Y)\|_2^2 \right] dt. \quad (8.13)$$

Show that the right-hand side equals $L(\theta) - L^*$, where $L(\theta)$ is the conditional flow matching loss (8.4) and $L^* = \int_0^1 \mathbb{E} \|\dot{I}_t - b_t(I_t, Y)\|_2^2 dt$ is the irreducible component of the loss independent of θ . *Hint.* Use the tower property and the fact that $b_t(I_t, Y) = \mathbb{E}[\dot{I}_t | I_t, Y]$ is the L^2 -projection, so $\mathbb{E} \|\hat{b}_t(I_t, Y, \theta) - \dot{I}_t\|^2 = \mathbb{E} \|\hat{b}_t(I_t, Y, \theta) - b_t(I_t, Y)\|^2 + \mathbb{E} \|b_t(I_t, Y) - \dot{I}_t\|^2$. Conclude that minimizing (8.4) directly controls the Y -averaged W_2^2 error between the learned and true posteriors.

- (c) *(Why pointwise accuracy is not guaranteed.)* Explain why the bound (8.13) does *not* imply pointwise accuracy $W_2^2(\hat{P}(\cdot | y), P(\cdot | y)) \leq \varepsilon$ for a specific y . Construct a simple example where the integrated CFM loss is small but the pointwise error is large for some y in a low-probability region of $P(y)$. *Hint.* Take a discrete $Y \in \{y_1, y_2\}$ with $P(Y = y_1) = 1 - \delta$ and $P(Y = y_2) = \delta$ for small δ . If the model has small error for y_1 and large error Δ for y_2 , the averaged error is $(1 - \delta) \cdot \text{small} + \delta \cdot \Delta^2$, which is small in expectation even when Δ is large. This is the precise sense in which conditional flow matching is *amortized*: rare observations may receive poor conditional samples.

2. **(Unification of conditional flow matching and guidance.)** This exercise shows that, at the level of exact solutions, Settings 1 and 2 of Section 8 are solving the

same equation: both produce a drift that is the unconditional flow matching formula (8.5) with the unconditional score $\nabla_x \log p_t(x)$ replaced by the conditional score $\nabla_x \log p_t(x | y)$. The two settings differ only in *how* they obtain the conditional score.

Throughout, use the linear interpolant $I_t = \alpha_t X_0 + \beta_t X_1$ with $X_0 \sim \mathcal{N}(0, I_d)$, $X_1 \sim \pi$, $Y | X_1 \sim P(\cdot | X_1)$. Let $p_t(x | y)$ denote the conditional density of I_t given $Y = y$.

- (a) (*Conditional analogue of Theorem 7.7.*) Conditional on $Y = y$ and $X_1 = x_1$, the interpolant $I_t | X_1 = x_1, Y = y$ has the same distribution as $I_t | X_1 = x_1$, namely $\mathcal{N}(\beta_t x_1, \alpha_t^2 I_d)$ (since $X_0 \perp Y$ given X_1). Apply Theorem 7.6 with $X = I_t | Y = y$ and $W = X_1 | Y = y$ to show that the conditional score satisfies

$$\nabla_x \log p_t(x | y) = -\frac{1}{\alpha_t^2} \mathbb{E}[\alpha_t X_0 | I_t = x, Y = y]. \quad (8.14)$$

- (b) (*Conditional velocity in terms of the conditional score.*) Following the derivation of (7.11), starting from $b_t(x, y) = \mathbb{E}[\dot{I}_t | I_t = x, Y = y] = \dot{\alpha}_t \mathbb{E}[X_0 | I_t = x, Y = y] + \dot{\beta}_t \mathbb{E}[X_1 | I_t = x, Y = y]$, and using (8.14) together with $x = \alpha_t \mathbb{E}[X_0 | I_t = x, Y = y] + \beta_t \mathbb{E}[X_1 | I_t = x, Y = y]$, derive

$$b_t(x, y) = \frac{\dot{\beta}_t}{\beta_t} x + \alpha_t^2 \left(\frac{\dot{\beta}_t}{\beta_t} - \frac{\dot{\alpha}_t}{\alpha_t} \right) \nabla_x \log p_t(x | y). \quad (8.15)$$

This is exactly (8.5) with the unconditional score $\nabla_x \log p_t(x)$ replaced by the conditional score $\nabla_x \log p_t(x | y)$.

- (c) (*Reconciling Settings 1 and 2.*) Combine (8.15) with the SDE family Theorem 7.8 (applied to the conditional marginals $p_t(\cdot | y)$): for any $g_t \geq 0$, the SDE

$$d\tilde{X}_t = (b_t(\tilde{X}_t, y) + g_t \nabla_x \log p_t(\tilde{X}_t | y)) dt + \sqrt{2g_t} dW_t, \quad \tilde{X}_0 \sim \mathcal{N}(0, I_d),$$

has $\text{Law}(\tilde{X}_t) = p_t(\cdot | y)$.

- (i) (*Setting 1 / CFM.*) Setting $g_t = 0$ recovers the conditional flow matching ODE (8.2). The conditional score $\nabla_x \log p_t(x | y)$ enters *implicitly* through $b_t(x, y)$, which is regressed onto $\dot{I}_t$ in the loss (8.4).
- (ii) (*Setting 2 / guidance.*) Apply Bayes' rule (8.7) to the conditional score in (8.15): $\nabla_x \log p_t(x | y) = \nabla_x \log p_t(x) + \nabla_x \log p_t(y | x)$. Substituting and rearranging:

$$b_t(x, y) = \underbrace{b_t(x)}_{\text{unconditional drift}} + \alpha_t^2 \left(\frac{\dot{\beta}_t}{\beta_t} - \frac{\dot{\alpha}_t}{\alpha_t} \right) \underbrace{\nabla_x \log p_t(y | x)}_{\text{likelihood score}}.$$

Compare with (8.11): this is exactly the guided drift, with the likelihood score appearing as the additive “guidance” correction to the unconditional drift.

Conclusion. Settings 1 and 2 produce identical *exact* drifts; they differ only in implementation:

- Setting 1 estimates the full conditional score $\nabla_x \log p_t(x \mid y)$ directly via regression onto $\dot{I}_t$ from paired data (X_1, Y) .
- Setting 2 decomposes the conditional score by Bayes into the (known) prior score plus the likelihood score, and approximates the latter using (8.9).

The two approaches incur different errors: Setting 1 has amortization error (Exercise 1); Setting 2 has the bias of (8.9). SMC-style methods correct the latter using the dynamical-IS perspective discussed in Section 8.3.

9 Stochastic Optimal Control for Conditional Sampling

9.1 Motivation

In Section 8.3 we approximated the conditional score $\nabla_x \log p_t(x | y)$ by Bayes' rule plus the guidance approximation (8.9), which introduces a bias. We now present a more principled approach via *stochastic optimal control* (SOC), which yields the *exact* conditional drift as the solution of an optimization problem (under certain assumptions) and in such a setting, (8.10) can be understood as a approximation of the control problem.

Setup. We have a pre-trained unconditional generative SDE

$$dX_t = b_t(X_t) dt + \sigma_t dW_t, \quad X_0 \sim \pi_0 = \mathcal{N}(0, I_d), \quad (9.1)$$

with $\text{Law}(X_1) = \pi_1$ (the data distribution π). Examples include the reverse SDE (6.4) (after time-reversal) and the SDE family (7.12). The *real goal* is to sample from the tilted distribution

$$\pi_1^R(x) \propto \pi_1(x) \exp(R(x)), \quad (9.2)$$

where $R : \mathbb{R}^d \rightarrow \mathbb{R}$ is a reward (or the negative log-likelihood $\log P(y | x)$ in Bayesian inference for a fixed observation y). By itself, (9.1) samples π_1 , not π_1^R .

Steering via a control. Add a control term $u_t : \mathbb{R}^d \rightarrow \mathbb{R}^d$ to the drift and consider the *controlled SDE*:

$$dX_t^u = (b_t(X_t^u) + \sigma_t u_t(X_t^u)) dt + \sigma_t dW_t, \quad X_0^u \sim \pi_0, \quad (9.3)$$

parameterized so that taking $u_t \equiv 0$ recovers (9.1). The question becomes: what control u_t steers (9.3) so that $\text{Law}(X_1^u) = \pi_1^R$? The guidance correction $f_t(x, y)$ in (8.10) is one approximation answer; SOC provides a principled framework.

9.2 The Stochastic Optimal Control Problem

Cost functional. Consider the SOC problem

$$\min_{u_t} \int_0^1 \frac{1}{2} \mathbb{E}[\|u_t(X_t^u)\|_2^2] dt - \mathbb{E}[R(X_1^u)] \quad \text{s.t. (9.3)}, \quad (9.4)$$

which combines a *running cost* (the L^2 control energy) with a *terminal reward* $R(X_1^u)$.

Why this cost? The running cost penalizes deviation from the uncontrolled process, keeping X_t^u close in distribution to X_t ; the terminal cost rewards trajectories ending in high- R regions. We will see in Section 9.5 that $\frac{1}{2} \int_0^1 \mathbb{E}\|u_t\|_2^2 dt$ is exactly $D_{\text{KL}}(\mathbb{P}^u \parallel \mathbb{P})$, the path-space KL divergence between the controlled and uncontrolled measures. The SOC problem (9.4) is therefore a KL-regularized maximization of $\mathbb{E}[R]$.

9.3 Value Function and Dynamic Programming

Value function. Define the cost-to-go starting from state x at time t :

$$V_t(x) := \min_u \mathbb{E} \left[\frac{1}{2} \int_t^1 \|u_s(X_s^u)\|_2^2 ds - R(X_1^u) \mid X_t^u = x \right], \quad (9.5)$$

so that $V_1(x) = -R(x)$ is the terminal condition.

Proposition 9.1 (Dynamic programming). *For $\Delta t > 0$,*

$$V_t(x) = \min_{u_t} \mathbb{E} \left[\frac{1}{2} \int_t^{t+\Delta t} \|u_s(X_s^u)\|_2^2 ds + V_{t+\Delta t}(X_{t+\Delta t}^u) \mid X_t^u = x \right]. \quad (9.6)$$

The proof is the standard splitting of the cost integral at $t + \Delta t$ and using the Markov property of the controlled SDE.

9.4 The Hamilton–Jacobi–Bellman Equation

Heuristic derivation. Send $\Delta t \rightarrow 0$ in (9.6). Expanding $V_{t+\Delta t}(X_{t+\Delta t}^u)$ via Itô's formula along the controlled SDE (9.3):

$$V_{t+\Delta t}(X_{t+\Delta t}^u) \approx V_t(x) + \partial_t V_t(x) \Delta t + \nabla V_t(x) \cdot \Delta x + \frac{1}{2} \sigma_t^2 \Delta V_t(x) \Delta t + O((\Delta t)^2),$$

where $\Delta x = (b_t(x) + \sigma_t u_t(x)) \Delta t + \sigma_t \Delta W_t$ and $\mathbb{E}[\Delta x] = (b_t + \sigma_t u_t) \Delta t$. Substituting into (9.6) and dividing by Δt :

$$0 = \min_{u_t} \left[\frac{1}{2} \|u_t(x)\|_2^2 + \partial_t V_t(x) + \nabla V_t(x) \cdot (b_t(x) + \sigma_t u_t(x)) + \frac{1}{2} \sigma_t^2 \Delta V_t(x) \right]. \quad (9.7)$$

The minimization over u_t is pointwise and quadratic, with minimizer

$$u_t^*(x) = -\sigma_t \nabla V_t(x). \quad (9.8)$$

Substituting (9.8) back into (9.7) eliminates the minimization and gives the *Hamilton–Jacobi–Bellman* (HJB) equation.

Proposition 9.2 (Optimal control and HJB). *The optimal control is $u_t^*(x) = -\sigma_t \nabla V_t(x)$, and the value function V_t satisfies the HJB equation*

$$\partial_t V_t(x) + b_t(x) \cdot \nabla V_t(x) - \frac{1}{2} \sigma_t^2 \|\nabla V_t(x)\|_2^2 + \frac{1}{2} \sigma_t^2 \Delta V_t(x) = 0, \quad V_1(x) = -R(x). \quad (9.9)$$

Remark 9.3. The HJB equation is a *nonlinear* PDE due to the $\|\nabla V_t\|^2$ term, reflecting the optimization over u . The Cole–Hopf transformation $V_t(x) = -\log \Phi_t(x)$ linearizes it: Φ_t then satisfies a linear (backward) heat equation, the Feynman–Kac representation of which connects Φ_t to the unconditional process (9.1). This linearization is the foundation of the SMC and twisted particle methods discussed in Section 8.3.

9.5 Characterizing the Distributions via Girsanov's Theorem

The remaining question is: what is the law of X_1^{u*} under the optimal control? We answer it via the Girsanov change of measure.

Path measures. Let $\mathbb{P}$ denote the path measure of the uncontrolled process $\{X_s : 0 \leq s \leq 1\}$ from (9.1), and $\mathbb{P}^u$ the path measure of the controlled process $\{X_s^u : 0 \leq s \leq 1\}$ from (9.3).

Discrete-time intuition. Discretize: $X_{t_{n+1}} = X_{t_n} + b_{t_n} \Delta t + \sigma_{t_n} (W_{t_{n+1}} - W_{t_n})$ with grid $0 = t_0 < \dots < t_N = 1$. The Markov factorization gives

$$\mathbb{P}(X) \approx p(X_{t_0}) \prod_k p(X_{t_k} | X_{t_{k-1}}), \quad \mathbb{P}^u(X) \approx p(X_{t_0}) \prod_k p^u(X_{t_k} | X_{t_{k-1}}),$$

where the conditional densities are Gaussian:

$$p(X_{t_k} | X_{t_{k-1}}) = \frac{1}{(2\pi\sigma_{t_{k-1}}^2 \Delta t)^{d/2}} \exp\left(-\frac{\|X_{t_k} - X_{t_{k-1}} - b_{t_{k-1}} \Delta t\|^2}{2\sigma_{t_{k-1}}^2 \Delta t}\right),$$

$$p^u(X_{t_k} | X_{t_{k-1}}) = \frac{1}{(2\pi\sigma_{t_{k-1}}^2 \Delta t)^{d/2}} \exp\left(-\frac{\|X_{t_k} - X_{t_{k-1}} - (b_{t_{k-1}} + \sigma_{t_{k-1}} u_{t_{k-1}}) \Delta t\|^2}{2\sigma_{t_{k-1}}^2 \Delta t}\right).$$

Computing the log-ratio and summing:

$$\log \frac{\mathbb{P}^u}{\mathbb{P}}(X) \approx \sum_k \left[-\frac{1}{2} \|u_{t_{k-1}}\|_2^2 \Delta t + \frac{1}{\sigma_{t_{k-1}}} (X_{t_k} - X_{t_{k-1}} - b_{t_{k-1}} \Delta t) \cdot u_{t_{k-1}} \right].$$

Continuous-time limit. Sending $\Delta t \rightarrow 0$ formally, the second sum in the discrete log-ratio becomes a path integral involving dX_t :

$$\log \frac{d\mathbb{P}^u}{d\mathbb{P}}(X) = -\frac{1}{2} \int_0^1 \|u_t(X_t)\|_2^2 dt + \int_0^1 \frac{1}{\sigma_t} (dX_t - b_t(X_t) dt) \cdot u_t(X_t). \quad (9.10)$$

This is the Radon–Nikodym density of $\mathbb{P}^u$ with respect to $\mathbb{P}$, evaluated on a path $X_{[0,1]}$. The expression (9.10) makes no reference to *which* measure the path X is sampled from; it is a function of the path. The interpretation is determined by the measure under which we take expectations.

Evaluating along the controlled process. We pick $X = X^u$ as the controlled process with law $\mathbb{P}^u$, satisfying (9.3). Then $dX_t^u - b_t(X_t^u) dt = \sigma_t u_t(X_t^u) dt + \sigma_t dW_t$, where W_t is a Brownian motion under $\mathbb{P}^u$. Substituting into (9.10):

$$\frac{1}{\sigma_t} (dX_t^u - b_t(X_t^u) dt) \cdot u_t(X_t^u) = \|u_t(X_t^u)\|_2^2 dt + u_t(X_t^u) \cdot dW_t,$$

so

$$\log \frac{d\mathbb{P}^u}{d\mathbb{P}}(X^u) = \frac{1}{2} \int_0^1 \|u_t(X_t^u)\|_2^2 dt + \int_0^1 u_t(X_t^u) \cdot dW_t. \quad (9.11)$$

Taking expectation under $\mathbb{P}^u$, the stochastic integral $\int_0^1 u_t \cdot dW_t$ has mean zero (Itô integral against a $\mathbb{P}^u$ -Brownian motion):

Theorem 9.4 (Girsanov / KL of path measures). *Under suitable integrability conditions (e.g., Novikov), the Radon–Nikodym density of $\mathbb{P}^u$ with respect to $\mathbb{P}$ is given by (9.10), and*

$$D_{\text{KL}}(\mathbb{P}^u \parallel \mathbb{P}) = \mathbb{E}^{\mathbb{P}^u} \left[\log \frac{d\mathbb{P}^u}{d\mathbb{P}}(X^u) \right] = \frac{1}{2} \int_0^1 \mathbb{E}^{\mathbb{P}^u} [\|u_t(X_t^u)\|_2^2] dt. \quad (9.12)$$

Remark 9.5 (Evaluating along the uncontrolled process). Alternatively, picking X as the uncontrolled process under $\mathbb{P}$ (so $dX_t - b_t(X_t) dt = \sigma_t dW_t$) gives

$$\log \frac{d\mathbb{P}^u}{d\mathbb{P}}(X) = -\frac{1}{2} \int_0^1 \|u_t(X_t)\|_2^2 dt + \int_0^1 u_t(X_t) \cdot dW_t,$$

with sign opposite to (9.11). This form is useful for importance-sampling computations: to estimate $\mathbb{E}^{\mathbb{P}^u}[\Phi(X)]$ by sampling from $\mathbb{P}$, reweight by $d\mathbb{P}^u/d\mathbb{P}$ evaluated along uncontrolled paths.

9.6 Connection to Sampling: KL-Control Duality

KL-control duality. Combining the SOC problem (9.4) with the path-KL formula (9.12):

$$\min_u [D_{\text{KL}}(\mathbb{P}^u \parallel \mathbb{P}) - \mathbb{E}^{\mathbb{P}^u}[R(X_1^u)]] = \min_u D_{\text{KL}}(\mathbb{P}^u \parallel \mathbb{P}^R) + \text{const}, \quad (9.13)$$

where $\mathbb{P}^R$ is the path measure characterized by the Radon–Nikodym density

$$\frac{d\mathbb{P}^R}{d\mathbb{P}}(X_{[0,1]}) = \frac{e^{R(X_1)}}{\mathbb{E}_{\mathbb{P}}[e^{R(X_1)}]}. \quad (9.14)$$

The terminal marginal of $\mathbb{P}^R$ is precisely the desired tilted target:

$$\text{Law}_{\mathbb{P}^R}(X_1) \propto \pi_1(x) e^{R(x)} = \pi_1^R(x). \quad (9.15)$$

Thus, SOC can be thought of as KL optimization in the path space, namely variational inference in the path space.

Subtlety: the initial-distribution constraint. However, the controlled SDE (9.3) *constrains* $X_0^u \sim \pi_0 = \mathcal{N}(0, I_d)$, and the SOC problem (9.4) therefore minimizes (9.13) *only over* path measures $\mathbb{P}^u$ with this fixed initial marginal. The path measure $\mathbb{P}^R$ in (9.14) *does not* have $X_0 \sim \mathcal{N}(0, I_d)$ in general: the tilting by $e^{R(X_1)}$ couples X_1 to X_0 through the joint $\mathbb{P}$, so its X_0 -marginal is

$$\text{Law}_{\mathbb{P}^R}(X_0) \propto \pi_0(x_0) \mathbb{E}_{\mathbb{P}}[e^{R(X_1)} \mid X_0 = x_0],$$

which differs from π_0 unless $\mathbb{E}_{\mathbb{P}}[e^{R(X_1)} \mid X_0]$ is constant in X_0 — typically not the case, since X_0 is informative about X_1 under the generative dynamics.

Consequence. The constrained SOC problem therefore does *not* produce $\mathbb{P}^u = \mathbb{P}^R$, and the law of $X_1^{u^*}$ under the constrained optimal control is in general *not exactly* π_1^R . This work of Domingo-Enrich, Drozdal, Karrer & Chen (2024) [10] show that exact sampling from π_1^R requires a special noise schedule (“memoryless” noise) so that X_0 becomes informationally decoupled from X_1 along the controlled path; only then does the constrained optimum coincide with $\mathbb{P}^R$ and yield $\text{Law}(X_1^{u^*}) = \pi_1^R$. Without this modification, the SOC solution provides a high-quality but slightly biased approximation to π_1^R .

Remark 9.6 (Memoryless noise schedule). The fix is to choose σ_t large enough so that the controlled process forgets X_0 before reaching $t = 1$. Concretely, when the unconditional SDE (9.1) is constructed so that $\mathbb{E}_{\mathbb{P}}[e^{R(X_1)} \mid X_0 = x_0]$ is independent of x_0 (the memoryless condition), the constrained KL-minimization attains $\mathbb{P}^u = \mathbb{P}^R$ and the SOC solution samples exactly from π_1^R . See [10] for the precise formulation.

We note that for score-based diffusion models based on time reversal of Langevin, the SDE automatically uses the memoryless schedule due to the ergodicity of the forward noising diffusion, which implies that the noise and target are independently coupled.

Connection to guidance. The optimal drift in (9.3) under u^* is

$$b_t(x) + \sigma_t u_t^*(x) = b_t(x) - \sigma_t^2 \nabla V_t(x),$$

which has the same affine-in-correction structure as the guided drift (8.11): the guidance term $f_t(x, y)$ plays the role of $-\sigma_t \nabla V_t(x)$ (up to the SDE-family prefactor), and the guidance approximation (8.9) corresponds to taking

$$V_t(x) \approx -\log p_1(y \mid \mathbb{E}[X_1 \mid I_t = x]).$$

9.7 Pathwise Gradient and the Adjoint Equation

There are two natural strategies for solving the SOC problem (9.4):

- (i) *Solve the optimality condition directly.* The HJB equation (9.9) is the optimality condition for the SOC problem; given V_t , the optimal control is recovered as $u_t^*(x) = -\sigma_t \nabla V_t(x)$. This requires solving a nonlinear PDE in $\mathbb{R}^d$.
- (ii) *Optimize the objective directly.* Parameterize the control as u_θ and minimize the SOC cost $J(\theta)$ by stochastic gradient descent. This requires computing $\nabla_\theta J(\theta)$ along simulated trajectories.

This subsection develops the second route. The key tool is the *adjoint equation*, which collapses backpropagation through a simulated trajectory into a single backward ODE and yields gradients of J with respect to either u (functional gradient) or θ (parameter gradient).

Roadmap.

- Section 9.7.1: derive the adjoint recursion in discrete time; this is just reverse-mode autodiff through Euler–Maruyama, transparent and intuitive.
- Section 9.7.2: take the formal continuum limit to obtain the pathwise adjoint ODE.
- Sections 9.7.3 and 9.7.4: use the adjoint to express the functional gradient $\delta J/\delta u$ and the parameter gradient $\nabla_\theta J$.
- Section 9.7.5: discuss *adjoint matching* [10], a recent regression-based simplification that uses a *lean* adjoint (dropping all u -dependent terms from the adjoint ODE to simplify it) and converts the SOC problem into a least-squares loss.

Generalized cost. Throughout this subsection we work with the slightly more general SOC cost

$$J(u) := \mathbb{E} \left[\int_0^1 \left(\frac{1}{2} \|u_t(X_t^u)\|_2^2 + f_t(X_t^u) \right) dt + g(X_1^u) \right], \quad (9.16)$$

allowing a state-dependent running cost f_t and terminal cost g . The setup of (9.4) is the special case $f \equiv 0$ and $g = -R$.

Notation. We use lowercase x for a generic state argument (so $u_t(x)$ is the control as a function of x and t) and capital X_t for the random state along a trajectory. Gradients $\nabla u_t(x)$, $\nabla b_t(x)$, $\nabla f_t(x)$, etc. denote derivatives with respect to x ; when evaluated at the current trajectory we write $\nabla u_t(X_t)$ etc.

9.7.1 Discrete Pathwise Adjoint

Discretize (9.3) on the grid $0 = t_0 < t_1 < \dots < t_N = 1$ with $\Delta t = 1/N$:

$$X_{k+1} = X_k + (b_k(X_k) + \sigma_k u_k(X_k)) \Delta t + \sigma_k \sqrt{\Delta t} \xi_k, \quad \xi_k \stackrel{\text{iid}}{\sim} \mathcal{N}(0, I_d), \quad (9.17)$$

with feedback control $u_k(\cdot) = u(\cdot, t_k)$ and shorthand $b_k = b_{t_k}$, $\sigma_k = \sigma_{t_k}$, $f_k = f_{t_k}$.

Pathwise adjoint. Freeze the noise sequence $\{\xi_k\}_{k=0}^{N-1}$. Along this fixed path, the trajectory $\{X_j\}_{j \geq k}$ is a deterministic function of the state X_k at step k via the recursion (9.17). Hence the future cost

$$\mathcal{C}_k(x) := \sum_{j=k}^{N-1} \left(\frac{1}{2} \|u_j(X_j(x))\|^2 + f_j(X_j(x)) \right) \Delta t + g(X_N(x))$$

is a deterministic function of the initial state x at step k , where $X_j(x)$ denotes the trajectory started from $X_k = x$ along the fixed noise sequence. Define the *pathwise adjoint* as its gradient evaluated at the actual state:

$$a_k := \nabla \mathcal{C}_k(X_k). \quad (9.18)$$

Equivalently, a_k is the sensitivity of the path-wise tail cost to a perturbation of X_k .

Backward recursion. Introduce the per-step running cost $\ell_t(x) := \frac{1}{2}\|u_t(x)\|^2 + f_t(x)$. Splitting off the k -th term and applying the chain rule,

$$\mathcal{C}_k(x) = \ell_k(x) \Delta t + \mathcal{C}_{k+1}(F_k(x)),$$

where $F_k(x) := x + (b_k(x) + \sigma_k u_k(x))\Delta t + \sigma_k \sqrt{\Delta t} \xi_k$ is the one-step Euler map. Differentiating in x and evaluating at X_k :

$$a_k = \nabla \ell_k(X_k) \Delta t + \nabla F_k(X_k)^\top a_{k+1}. \quad (9.19)$$

The one-step Jacobian, from F_k , is

$$\nabla F_k(X_k) = I + (\nabla b_k(X_k) + \sigma_k \nabla u_k(X_k)) \Delta t, \quad (9.20)$$

where the noise term $\sigma_k \sqrt{\Delta t} \xi_k$ contributes *nothing* because σ_k is state-independent. Substituting into (9.19):

$$\boxed{a_k = \nabla \ell_k(X_k) \Delta t + (I + (\nabla b_k(X_k) + \sigma_k \nabla u_k(X_k)) \Delta t)^\top a_{k+1}, \quad a_N = \nabla g(X_N).} \quad (9.21)$$

This recursion is precisely reverse-mode automatic differentiation through the Euler–Maruyama scheme (9.17).

9.7.2 Continuous Pathwise Adjoint ODE

Rearranging (9.21) as

$$\frac{a_{k+1} - a_k}{\Delta t} = -[(\nabla b_k(X_k) + \sigma_k \nabla u_k(X_k))^\top a_{k+1} + \nabla \ell_k(X_k)] + O(\Delta t),$$

and sending $\Delta t \rightarrow 0$:

Proposition 9.7 (Pathwise adjoint ODE). *Along each Brownian path of (9.3), the pathwise adjoint*

$$a_t := \nabla_x \left[\int_t^1 \ell_s(X_s(x)) ds + g(X_1(x)) \right] \Big|_{x=X_t}, \quad \ell_s(x) := \frac{1}{2}\|u_s(x)\|^2 + f_s(x),$$

where $X_s(x)$ denotes the trajectory started from $X_t = x$ along the fixed Brownian path, satisfies the linear backward ODE

$$\frac{da_t}{dt} = -[(\nabla b_t(X_t) + \sigma_t \nabla u_t(X_t))^\top a_t + \nabla \ell_t(X_t)], \quad a_1 = \nabla g(X_1). \quad (9.22)$$

Remark 9.8 (Pathwise vs. adapted: BSDE form). By construction, a_t is $\mathcal{F}_1$ -measurable but *not* $\mathcal{F}_t$ -adapted: it depends on the entire future Brownian path through the trajectory $\{X_s\}_{s \in [t, 1]}$. Its $\mathcal{F}_t$ -projection

$$\hat{a}_t := \mathbb{E}[a_t \mid \mathcal{F}_t] = \nabla_x J(u; X_t, t), \quad (9.23)$$

where $J(u; x, t)$ is the cost functional from (9.5) with the min removed, is adapted (Markov property + interchange of ∇ and $\mathbb{E}$). Applying Itô's formula to $\nabla_x J(u; X_t, t)$ along the controlled SDE, using the linear Feynman–Kac PDE for J , yields the *backward stochastic differential equation* (BSDE)

$$d\hat{a}_t = -[(\nabla b_t(X_t) + \sigma_t \nabla u_t(X_t))^\top \hat{a}_t + \nabla \ell_t(X_t)] dt + Z_t dW_t, \quad \hat{a}_1 = \nabla g(X_1), \quad (9.24)$$

where the martingale coefficient Z_t is determined by the martingale representation theorem. The drift agrees with the pathwise ODE (9.22); the $Z_t dW_t$ piece is the price of forcing adaptedness. The pathwise ODE (9.22) is what one *implements* via backpropagation through a single simulated trajectory; the BSDE (9.24) is the form for theoretical analysis (FBSDE methods).

9.7.3 Functional Gradient

We use the pathwise adjoint to compute the gradient of the SOC cost (9.16) with respect to u as a function. Perturb $u \rightarrow u + \varepsilon \delta u$ and let X_t^ε denote the controlled trajectory under the perturbed control. The first-order trajectory variation $\delta X_t := \partial_\varepsilon X_t^\varepsilon|_{\varepsilon=0}$ satisfies the linearized SDE

$$d(\delta X_t) = (\nabla b_t(X_t) + \sigma_t \nabla u_t(X_t)) \delta X_t dt + \sigma_t \delta u_t(X_t) dt, \quad \delta X_0 = 0, \quad (9.25)$$

with no Itô term, since σ_t is state-independent. Differentiating (9.16) at $\varepsilon = 0$ (recall $\ell_t(x) = \frac{1}{2} \|u_t(x)\|^2 + f_t(x)$):

$$\left. \frac{dJ}{d\varepsilon} \right|_0 = \mathbb{E} \left[\int_0^1 \langle u_t(X_t), \delta u_t(X_t) \rangle dt + \int_0^1 \nabla \ell_t(X_t)^\top \delta X_t dt + \nabla g(X_1)^\top \delta X_1 \right]. \quad (9.26)$$

The first integral is the *direct* variation through u_t at fixed trajectory; the remaining two are the *indirect* variation through δX_t .

Adjoint identity. The dynamics (9.22) are designed so that the indirect contribution collapses. Compute, using (9.22) and (9.25):

$$\begin{aligned} \frac{d}{dt}(a_t^\top \delta X_t) &= \dot{a}_t^\top \delta X_t + a_t^\top \dot{\delta X}_t \\ &= -[(\nabla b_t(X_t) + \sigma_t \nabla u_t(X_t))^\top a_t + \nabla \ell_t(X_t)]^\top \delta X_t \\ &\quad + a_t^\top [(\nabla b_t(X_t) + \sigma_t \nabla u_t(X_t)) \delta X_t + \sigma_t \delta u_t(X_t)] \\ &= -\nabla \ell_t(X_t)^\top \delta X_t + a_t^\top \sigma_t \delta u_t(X_t), \end{aligned}$$

where the $(\nabla b + \sigma \nabla u)$ -terms cancel between the adjoint drift and the linearized state. Integrating from 0 to 1 with $\delta X_0 = 0$ and $a_1 = \nabla g(X_1)$:

$$\nabla g(X_1)^\top \delta X_1 + \int_0^1 \nabla \ell_t(X_t)^\top \delta X_t dt = \int_0^1 \langle \sigma_t^\top a_t, \delta u_t(X_t) \rangle dt. \quad (9.27)$$

The left side is the indirect contribution in (9.26). Substituting:

Proposition 9.9 (Functional gradient). *The first variation of J along an admissible perturbation δu is*

$$\left. \frac{dJ}{d\varepsilon} \right|_0 = \mathbb{E} \left[\int_0^1 \langle u_t(X_t) + \sigma_t^\top a_t, \delta u_t(X_t) \rangle dt \right], \quad (9.28)$$

where a_t is the pathwise adjoint (9.22). Equivalently, in terms of the law ρ_t^u of X_t^u ,

$$\frac{\delta J}{\delta u}(x, t) = \rho_t^u(x) [u_t(x) + \sigma_t^\top \mathbb{E}[a_t \mid X_t = x]].$$

Remark 9.10 (Recovering HJB optimality). The stationarity condition $\delta J / \delta u = 0$ reads

$$u_t^*(x) = -\sigma_t^\top \mathbb{E}[a_t \mid X_t = x] = -\sigma_t^\top \nabla_x J(u^*; x, t) = -\sigma_t^\top \nabla V_t(x),$$

using (9.23) and $\nabla J(u^*; \cdot) = \nabla V$ at the optimum. This recovers the HJB optimality formula (9.8) as a function-space stationarity condition, complementing the dynamic programming derivation of Theorem 9.2.

9.7.4 Parameter Gradient

For numerical optimization, parameterize the control by a neural network $u_t(x, \theta)$ with parameters $\theta \in \mathbb{R}^p$. The SOC cost becomes $J(\theta) := J(u_t(\cdot, \theta))$, and we want $\nabla_\theta J$. The derivation parallels the functional-gradient case.

Linearized trajectory. Perturb $\theta \rightarrow \theta + \varepsilon \delta \theta$ and let X_t^ε be the trajectory under the perturbed control $u_t(\cdot, \theta + \varepsilon \delta \theta)$. Linearizing (9.3) in ε , the trajectory variation $\delta X_t := \partial_\varepsilon X_t^\varepsilon|_{\varepsilon=0}$ satisfies

$$d(\delta X_t) = (\nabla b_t(X_t) + \sigma_t^\top \nabla_x u_t(X_t, \theta)) \delta X_t dt + \sigma_t^\top \partial_\theta u_t(X_t, \theta) \delta \theta dt, \quad \delta X_0 = 0, \quad (9.29)$$

where $\nabla_x u_t(x, \theta) \in \mathbb{R}^{d \times d}$ is the state Jacobian of $u_t(\cdot, \theta)$ at fixed θ , and $\partial_\theta u_t(x, \theta) \in \mathbb{R}^{d \times p}$ is the parameter Jacobian at fixed (x, t) .

Cost variation. The control cost $\frac{1}{2} \|u_t(X_t^\varepsilon, \theta)\|^2$ depends on ε both through X_t^ε and through $u_t(\cdot, \theta)$ explicitly. By the chain rule, evaluating at X_t ,

$$\left. \frac{d}{d\varepsilon} \right|_0 \left[\frac{1}{2} \|u_t(X_t^\varepsilon, \theta)\|^2 \right] = u_t(X_t, \theta)^\top [\nabla_x u_t(X_t, \theta) \delta X_t + \partial_\theta u_t(X_t, \theta) \delta \theta],$$

where the first piece equals $\nabla_x (\frac{1}{2} \|u_t(x, \theta)\|^2)|_{x=X_t}^\top \delta X_t$. With $\ell_t(x) := \frac{1}{2} \|u_t(x, \theta)\|^2 + f_t(x)$,

$$\left. \frac{dJ}{d\varepsilon} \right|_0 = \underbrace{\mathbb{E} \left[\int_0^1 \nabla \ell_t(X_t)^\top \delta X_t dt + \nabla g(X_1)^\top \delta X_1 \right]}_{\text{indirect (through trajectory)}} + \underbrace{\mathbb{E} \left[\int_0^1 u_t(X_t, \theta)^\top \partial_\theta u_t(X_t, \theta) \delta \theta dt \right]}_{\text{direct (through } u_t)}. \quad (9.30)$$

The indirect term can be handled by the same identity as in Theorem 9.9, with δX_t now driven by $\delta\theta$ via (9.29) instead of by δu via (9.25). Compute, using (9.22) and (9.29):

$$\frac{d}{dt}(a_t^\top \delta X_t) = -\nabla \ell_t(X_t)^\top \delta X_t + a_t^\top \sigma_t \partial_\theta u_t(X_t, \theta) \delta\theta,$$

once again with the $(\nabla b_t + \sigma_t \nabla_x u_t)$ -terms cancelling. Integrating with $\delta X_0 = 0$ and $a_1 = \nabla g(X_1)$:

$$\nabla g(X_1)^\top \delta X_1 + \int_0^1 \nabla \ell_t(X_t)^\top \delta X_t dt = \int_0^1 a_t^\top \sigma_t \partial_\theta u_t(X_t, \theta) \delta\theta dt,$$

which is exactly the indirect term in (9.30). Substituting and reading off the coefficient of $\delta\theta$:

Proposition 9.11 (Parameter gradient). *Let $u_t(x, \theta)$ be a control parameterized smoothly by $\theta \in \mathbb{R}^p$. Then*

$$\boxed{\nabla_\theta J(\theta) = \mathbb{E} \left[\int_0^1 \partial_\theta u_t(X_t, \theta)^\top (u_t(X_t, \theta) + \sigma_t^\top a_t) dt \right]}, \quad (9.31)$$

where X_t is the trajectory generated under the current control $u_t(\cdot, \theta)$ and a_t is the pathwise adjoint (9.22) along that trajectory.

Direct vs. indirect contributions. Splitting (9.31):

$$\nabla_\theta J(\theta) = \frac{1}{2} \mathbb{E} \left[\int_0^1 \frac{\partial}{\partial \theta} \|u_t(X_t, \theta)\|^2 dt \right] + \mathbb{E} \left[\int_0^1 \partial_\theta u_t(X_t, \theta)^\top \sigma_t^\top a_t dt \right],$$

where the partial $\partial/\partial\theta$ in the first term holds X_t fixed. The first piece is the direct dependence on θ through the running cost $\frac{1}{2}\|u_t\|^2$; the second is the indirect dependence through the trajectory. The adjoint state a_t collapses the entire trajectory backpropagation into a single time-integrated coupling $\sigma_t^\top a_t$ at each t , so the cost of (9.31) is one forward simulation plus one backward ODE per gradient evaluation, independent of the parameter dimension p .

9.7.5 Adjoint Matching as Regression

From parameter gradient to regression loss. The functional form of (9.31) suggests training $u_t(\cdot, \theta)$ by minimizing the quadratic regression loss

$$\mathcal{L}_{\text{AM}}(\theta) := \frac{1}{2} \mathbb{E} \left[\int_0^1 \|u_t(\bar{X}_t, \theta) + \sigma_t^\top \bar{a}_t\|^2 dt \right], \quad (9.32)$$

where the bars on $\bar{X}_t$ and $\bar{a}_t$ indicate that, although both are simulated using the current value of θ on each iteration, they are treated as *constants* in θ when differentiating: any implicit dependence through the trajectory and adjoint is dropped, and only the

explicit θ -dependence in $u_t(\bar{X}_t, \theta)$ contributes to $\nabla_\theta \mathcal{L}_{\text{AM}}$. In an autodiff framework this is implemented by a `stop_gradient` (or `.detach()`) call applied to $\bar{X}_t$ and $\bar{a}_t$.

Differentiating (9.32) under this convention gives

$$\nabla_\theta \mathcal{L}_{\text{AM}}(\theta) = \mathbb{E} \left[\int_0^1 \partial_\theta u_t(\bar{X}_t, \theta)^\top (u_t(\bar{X}_t, \theta) + \sigma_t^\top \bar{a}_t) dt \right] = \nabla_\theta J(\theta),$$

which is *exactly* (9.31): the barred quantities $\bar{X}_t$ and $\bar{a}_t$ have the same numerical values as the θ -dependent X_t and a_t (they are the same simulated trajectory and adjoint, just flagged as constants for autodiff), so the integrands agree pointwise. The reason stopping gradient through $\bar{X}_t$ and $\bar{a}_t$ is correct, rather than throwing away information, is that the chain-rule contribution from the implicit θ -dependence of the trajectory has *already been absorbed* into the $\sigma_t^\top a_t$ term by the adjoint identity (9.27). Differentiating through $\bar{X}_t$ again would double-count that contribution.

(9.32) therefore provides a *plain least-squares implementation* of the parameter gradient: at each step, simulate the forward and backward passes to obtain $(\bar{X}_t, \bar{a}_t)$, treat them as data, and update θ by stochastic gradient descent on the regression of $u_t(\bar{X}_t, \theta)$ onto the target $-\sigma_t^\top \bar{a}_t$.

Lean adjoint. Computing a_t via (9.22) requires gradients of the control through both (i) the drift Jacobian $\sigma_t^\top \nabla_x u_t(\bar{X}_t, \theta)$ and (ii) the running cost $\nabla(\frac{1}{2} \|u_t(x, \theta)\|^2)|_{x=\bar{X}_t} = \nabla_x u_t(\bar{X}_t, \theta)^\top u_t(\bar{X}_t, \theta)$. Both involve the state Jacobian of $u_t(\cdot, \theta)$, a second-order quantity when $u_t(\cdot, \theta)$ is a neural network of x . The *lean adjoint* [10] $\tilde{a}_t$ drops all u -dependent contributions to the adjoint dynamics — both the $\sigma_t^\top \nabla_x u_t(\bar{X}_t, \theta)$ term in the linearized drift and the $\nabla(\frac{1}{2} \|u_t(\cdot, \theta)\|^2)$ term in the running cost — while still running along the (frozen) controlled trajectory $\bar{X}_t$:

$$\frac{d\tilde{a}_t}{dt} = -[\nabla b_t(\bar{X}_t)^\top \tilde{a}_t + \nabla f_t(\bar{X}_t)], \quad \tilde{a}_1 = \nabla g(\bar{X}_1). \quad (9.33)$$

The lean dynamics depend on the control only through the trajectory $\bar{X}_t$; the right-hand side itself contains no u . Replacing $\bar{a}_t$ with $\tilde{a}_t$ in (9.32) preserves u^* as a critical point [10]. The savings are substantial: each adjoint step requires only ∇b_t (a fixed term from the base SDE).

Week 6 Exercises

1. **(HJB via Lagrange multipliers.)** Derive the HJB equation (9.9) by treating the Fokker–Planck equation as a constraint.

Reformulate the SOC problem (9.4) as

$$\min_{u, \rho^u} \int_0^1 \int \frac{1}{2} \|u_t\|_2^2 \rho_t^u dx dt - \int R(x) \rho_1^u(x) dx$$

subject to the FP equation $\partial_t \rho_t^u + \nabla \cdot (\rho_t^u (b_t + \sigma_t u_t)) = \frac{1}{2} \sigma_t^2 \Delta \rho_t^u$ with $\rho_0^u = \pi_0$. Introduce a Lagrange multiplier $\lambda_t(x)$ for this constraint and form the Lagrangian $\mathcal{L}[u, \rho^u, \lambda]$.

By integration by parts in t and x , write $\mathcal{L}$ in the form

$$\mathcal{L} = \int_0^1 \int \rho_t^u H_t(x, u_t, \lambda_t) dx dt + \text{boundary terms},$$

where H_t is a pointwise function. Show that:

- (a) Stationarity in u_t gives $u_t^*(x) = \sigma_t \nabla \lambda_t(x)$.
- (b) Stationarity in ρ_t^u for $t \in (0, 1)$, after substituting u_t^* , gives a backward PDE for λ_t .
- (c) The boundary term in $t = 1$ gives the terminal condition $\lambda_1 = R$.

Setting $V_t = -\lambda_t$, verify that the equations from (a)–(c) are exactly $u_t^* = -\sigma_t \nabla V_t$ and the HJB equation (9.9). This identifies the value function as the Lagrange multiplier of the FP constraint (up to sign).

2. **(Connection to entropy-regularized RL.)** Show that the SOC problem (9.4) is the continuous-time limit of an entropy-regularized RL problem.

Background on RL. Reinforcement learning studies an agent interacting with an environment over discrete time steps $n = 0, 1, \dots, N$. At step n the agent observes a *state* s_n , samples an *action* $a_n \sim \pi_n(\cdot | s_n)$ from a *policy* π_n , and the environment transitions to a new state s_{n+1} (in continuous-time control problems, the transition kernel is determined by the dynamics of the system and is not chosen by the agent). The agent receives a per-step reward $r(s_n, a_n)$ and a terminal reward $r_T(s_N)$. The goal is to choose $\pi = (\pi_0, \dots, \pi_{N-1})$ to maximize $\mathbb{E}_\pi[\sum_n r(s_n, a_n) + r_T(s_N)]$.

In *entropy-regularized* (or KL-regularized) RL, one fixes a *reference policy* $\pi_n^{\text{ref}}(\cdot | s)$ and penalizes deviation from it via a KL term:

$$\max_{\pi} \mathbb{E}_{\pi} \left[r_T(s_N) + \sum_{n=0}^{N-1} \left(r(s_n, a_n) - D_{\text{KL}} \left(\pi_n(\cdot | s_n) \parallel \pi_n^{\text{ref}}(\cdot | s_n) \right) \right) \right]. \quad (9.34)$$

This formulation is standard in robotics (where π^{ref} is a safe baseline) and in fine-tuning of language models, where π^{ref} is a pretrained model and r_T is a learned reward (the standard RLHF setup).

Discretization of the SOC problem. Discretize the controlled SDE (9.3) on a grid $0 = t_0 < \dots < t_N = 1$ with step Δt . At each step, the controlled and uncontrolled transition kernels are Gaussian:

$$q^u(x' | x) = \mathcal{N}(x'; x + (b_t + \sigma_t u_t) \Delta t, \sigma_t^2 \Delta t I_d), \quad q(x' | x) = \mathcal{N}(x'; x + b_t \Delta t, \sigma_t^2 \Delta t I_d).$$

View q^u as a policy and q as a reference policy.

- (a) Show that $D_{\text{KL}}(q^u(\cdot | x) \| q(\cdot | x)) = \frac{1}{2} \|u_t(x)\|_2^2 \Delta t$.
- (b) Conclude that the SOC running cost $\int_0^1 \frac{1}{2} \mathbb{E} \|u_t\|_2^2 dt$ is the continuous-time limit of $\sum_n \mathbb{E}[D_{\text{KL}}(q^u(\cdot | X_{t_n}) \| q(\cdot | X_{t_n}))]$ — a sum of per-step KL penalties between the policy and reference, i.e. the entropy-regularization term in (9.34).
- (c) Compare the SOC objective (9.4) (a *minimization*) with the RL objective (9.34) (a *maximization*). Multiply the SOC objective by -1 to convert it into a maximization:

$$\max_u \mathbb{E}[R(X_1^u)] - \int_0^1 \frac{1}{2} \mathbb{E} \|u_t\|_2^2 dt.$$

Using (b), the running cost is the continuous-time limit of $\sum_n \mathbb{E}[D_{\text{KL}}(q^u \| q)]$. Match this term-by-term against (9.34) to obtain the dictionary

$$s_n \leftrightarrow X_{t_n}, \quad \pi_n \leftrightarrow q^u, \quad \pi_n^{\text{ref}} \leftrightarrow q, \quad r_T \leftrightarrow R, \quad r \leftrightarrow 0.$$

The per-step reward r vanishes because the SOC running cost contains only the KL term and no reward term inside the time integral; the entire reward signal is concentrated at the terminal time $t = 1$ via $R(X_1^u)$.

The conclusion is that fine-tuning a generative model to a tilted target $\pi_1 \exp(R)$ via SOC is structurally the same problem as fine-tuning a reference policy with KL-regularized RL (the standard RLHF setup), and methods developed for one translate directly to the other.

3. **(REINFORCE / score-function gradient.)** The pathwise adjoint of Section 9.7 requires backpropagating through the simulated SDE, which in particular needs ∇g and ∇f_t — the gradients of the terminal and running cost in x . In many applications (e.g. reward fine-tuning where g is a non-differentiable preference score, or a black-box simulator output) those gradients are unavailable. The *score-function* or *REINFORCE* estimator computes $\nabla_{\theta} J$ using only function evaluations of the cost, by moving the gradient onto the sampling density.

(a) REINFORCE in basic RL. For a single-step decision problem with action $a \sim \pi_\theta(\cdot)$ and reward $r(a)$, derive the score-function identity

$$\nabla_\theta \mathbb{E}_{a \sim \pi_\theta}[r(a)] = \mathbb{E}_{a \sim \pi_\theta}[r(a) \nabla_\theta \log \pi_\theta(a)], \quad (9.35)$$

by writing the expectation as an integral and differentiating under the integral sign. Observe that r enters only through evaluations, not through its gradient. For a multi-step MDP with trajectory $\tau = (s_0, a_0, s_1, a_1, \dots, s_T)$, joint density $p_\theta(\tau) = p(s_0) \prod_n \pi_\theta(a_n | s_n) p(s_{n+1} | s_n, a_n)$, and total reward $R(\tau) = \sum_n r(s_n, a_n)$, show that

$$\nabla_\theta \mathbb{E}_{\tau \sim p_\theta}[R(\tau)] = \mathbb{E}_{\tau \sim p_\theta}\left[R(\tau) \sum_n \nabla_\theta \log \pi_\theta(a_n | s_n)\right],$$

where the transition density $p(s_{n+1} | s_n, a_n)$ contributes nothing because it does not depend on θ .

(b) Discrete-time SOC via the score trick. Apply REINFORCE to the discretized SOC problem on the grid $0 = t_0 < \dots < t_N = 1$ with the Euler–Maruyama scheme (9.17). Show that the joint trajectory density is

$$p_\theta(X_{0:N}) = \pi_0(X_0) \prod_{k=0}^{N-1} \mathcal{N}(X_{k+1}; X_k + (b_k(X_k) + \sigma_k u_{t_k}(X_k, \theta)) \Delta t, \sigma_k^2 \Delta t I),$$

and compute that the per-step score is

$$\nabla_\theta \log p_\theta(X_{k+1} | X_k) = \partial_\theta u_{t_k}(X_k, \theta)^\top \frac{X_{k+1} - X_k - (b_k(X_k) + \sigma_k u_{t_k}(X_k, \theta)) \Delta t}{\sigma_k} \quad (9.36)$$

which can be written as $\partial_\theta u_{t_k}(X_k, \theta)^\top \sqrt{\Delta t} \xi_k$.

Hence, separating the direct dependence of the cost on θ from the indirect dependence through the law of $X_{0:N}$,

$$\nabla_\theta J(\theta) = \mathbb{E} \left[\sum_{k=0}^{N-1} \partial_\theta u_{t_k}(X_k, \theta)^\top u_{t_k}(X_k, \theta) \Delta t + \Phi(X_{0:N}) \sum_{k=0}^{N-1} \partial_\theta u_{t_k}(X_k, \theta)^\top \sqrt{\Delta t} \xi_k \right], \quad (9.37)$$

where $\Phi(X_{0:N}) := \sum_k (\frac{1}{2} \|u_k\|^2 + f_k) \Delta t + g(X_N)$ is the path cost. Note that Φ enters only through evaluations, never through gradients.

(c) Continuous limit. Send $\Delta t \rightarrow 0$ in (9.37). Show that the score becomes a stochastic integral,

$$\sum_{k=0}^{N-1} \partial_\theta u_{t_k}(X_k, \theta)^\top \sqrt{\Delta t} \xi_k \longrightarrow \int_0^1 \partial_\theta u_t(X_t, \theta)^\top dW_t,$$

yielding the *continuous-time REINFORCE estimator*

$$\boxed{\nabla_\theta J(\theta) = \mathbb{E} \left[\int_0^1 \partial_\theta u_t(X_t, \theta)^\top u_t(X_t, \theta) dt + \Phi(X) \int_0^1 \partial_\theta u_t(X_t, \theta)^\top dW_t \right]}, \quad (9.38)$$

where $\Phi(X) = \int_0^1 \ell_t(X_t) dt + g(X_1)$. Compare with the pathwise adjoint formula (9.31): the direct term agrees, while the $\sigma_t^\top a_t$ piece, which captures the trajectory-wise sensitivity through backward simulation, is here replaced by the return $\Phi(X)$ multiplied by the (forward-only) score $\int \partial_\theta u_t(X_t, \theta)^\top dW_t$. In particular, computing (9.38) requires no backward pass and no gradient of g or f_t ; only their values along the trajectory are needed.

(d) Girsanov derivation. Re-derive (9.38) directly in continuous time using Girsanov's theorem. Let $\mathbb{P}$ be the law of the uncontrolled base process $dX_t = b_t(X_t) dt + \sigma_t dW_t^\mathbb{P}$ and let $\mathbb{P}^\theta$ be the law of the controlled process. The score for the parametric family $\{\mathbb{P}^\theta\}$ relative to the fixed reference $\mathbb{P}$ is

$$\nabla_\theta \log p_\theta(X) = \nabla_\theta \log \frac{d\mathbb{P}^\theta}{d\mathbb{P}}(X),$$

which is a path functional, well-defined for any path X .

We compute this score for a *generic* path X , then specialize to $X = X^\theta$. From (9.10), the log Radon–Nikodym density expressed as a path functional is

$$\log \frac{d\mathbb{P}^\theta}{d\mathbb{P}}(X) = \int_0^1 u_t(X_t, \theta)^\top \sigma_t^{-1} (dX_t - b_t(X_t) dt) - \frac{1}{2} \int_0^1 \|u_t(X_t, \theta)\|^2 dt,$$

where $\sigma_t^{-1}(dX_t - b_t dt)$ denotes the path-data expression for the driving noise (and reduces to $dW_t^\mathbb{P}$ under $\mathbb{P}$ or to $u_t(X_t, \theta) dt + dW_t^{\mathbb{P}^\theta}$ under $\mathbb{P}^\theta$). Holding the path X fixed and differentiating in θ :

$$\nabla_\theta \log \frac{d\mathbb{P}^\theta}{d\mathbb{P}}(X) = \int_0^1 \partial_\theta u_t(X_t, \theta)^\top \sigma_t^{-1} (dX_t - b_t dt) - \int_0^1 \partial_\theta u_t(X_t, \theta)^\top u_t(X_t, \theta) dt.$$

Now specialize to the controlled path: $X_t = X_t^\theta$ satisfies $dX_t^\theta - b_t dt = \sigma_t u_t(X_t^\theta, \theta) dt + \sigma_t dW_t^{\mathbb{P}^\theta}$, so $\sigma_t^{-1}(dX_t^\theta - b_t dt) = u_t(X_t^\theta, \theta) dt + dW_t^{\mathbb{P}^\theta}$. Substituting, the cross-terms $\int \partial_\theta u_t(X_t^\theta, \theta)^\top u_t(X_t^\theta, \theta) dt$ cancel and

$$\nabla_\theta \log \frac{d\mathbb{P}^\theta}{d\mathbb{P}} \Big|_{X^\theta} = \int_0^1 \partial_\theta u_t(X_t^\theta, \theta)^\top dW_t^{\mathbb{P}^\theta}. \quad (9.39)$$

This is a martingale, hence has zero mean under $\mathbb{P}^\theta$, as a score should.

The standard score-function identity then yields, for any path functional Φ_θ depending on θ both directly (through the integrand) and indirectly (through the law of X),

$$\nabla_\theta \mathbb{E}^{\mathbb{P}^\theta}[\Phi_\theta] = \mathbb{E}^{\mathbb{P}^\theta}[\nabla_\theta \Phi_\theta] + \mathbb{E}^{\mathbb{P}^\theta}[\Phi_\theta \cdot \nabla_\theta \log p_\theta(X^\theta)].$$

Applying this with $\Phi_\theta = \int_0^1 \ell_t(X_t) dt + g(X_1)$, using $\nabla_\theta \Phi_\theta = \int_0^1 \partial_\theta u_t(X_t, \theta)^\top u_t(X_t, \theta) dt$ (only $\frac{1}{2} \|u_t(X_t, \theta)\|^2$ contributes a direct gradient), and substituting (9.39) recovers (9.38).

(e) **Variance.** Argue heuristically why the REINFORCE estimator (9.38) typically has higher variance than the pathwise adjoint estimator (9.31): in REINFORCE the *entire* return $\Phi(X)$ multiplies a stochastic integral, whereas in the pathwise adjoint the local sensitivity $\sigma_t^\top a_t$ isolates the contribution of each time slice. The trade-off between score-function and pathwise gradient estimators is the central design choice in policy-gradient reinforcement learning.

10 Sampling as Optimization in Probability Space

The SOC formulation of Section 9 cast generative sampling as an optimization problem on a space of *path measures*: minimize a KL divergence subject to dynamical constraints. The same idea applies more broadly. Most modern sampling algorithms can be understood as instances of optimization on the space $\mathcal{P}(\mathbb{R}^d)$ of probability measures, with the target π playing the role of the minimizer of an energy functional. The following dictionary will guide the narrative in this section:

Optimization on $\mathbb{R}^d$	Sampling on $\mathcal{P}(\mathbb{R}^d)$
gradient descent	Langevin dynamics
heavy-ball / Nesterov momentum	underdamped Langevin
preconditioned methods (Newton, natural gradient)	preconditioned samplers (Riemannian Langevin/HMC, ensemble)

This section walks through each row, with concrete examples to build intuition. We do not aim for full proofs — many of the rate results require nontrivial technical work — but we will state the key formulas, work through a few revealing special cases (mostly involving Gaussians, where everything is computable), and explain why each algorithm earns its place in the table.

10.1 Langevin as KL Minimization

We begin on the optimization side, recall the gradient-flow viewpoint together with its proximal and preconditioned variants, and then transport the entire picture to probability space via the Wasserstein metric. The end result is a clean identification of Langevin dynamics as Wasserstein gradient flow of KL divergence, with the log-Sobolev inequality playing the role of the Polyak–Łojasiewicz condition.

Optimization on $\mathbb{R}^d$ and gradient flows

Let $V : \mathbb{R}^d \rightarrow \mathbb{R}$ be smooth and (for concreteness) α -strongly convex, with unique minimizer x^* . The simplest discrete iteration,

$$x_{k+1} = x_k - h \nabla V(x_k), \quad (10.1)$$

is the *forward-Euler* discretization of the gradient-flow ODE

$$\dot{x}_t = -\nabla V(x_t). \quad (10.2)$$

Under α -strong convexity, $\frac{d}{dt} \|x_t - x^*\|^2 = 2(x_t - x^*) \cdot (-\nabla V(x_t)) \leq -2\alpha \|x_t - x^*\|^2$, so by Grönwall $\|x_t - x^*\|^2 \leq e^{-2\alpha t} \|x_0 - x^*\|^2$.

A second route to the same flow is the *proximal* or backward-Euler scheme

$$x_{k+1} = \arg \min_{y \in \mathbb{R}^d} \left[V(y) + \frac{1}{2h} \|y - x_k\|^2 \right]. \quad (10.3)$$

The first-order condition gives the implicit update $x_{k+1} = x_k - h \nabla V(x_{k+1})$, which is backward Euler for (10.2). Both (10.1) and (10.3) limit to the gradient flow (10.2) as $h \rightarrow 0$. The difference matters when $h > 0$: forward Euler is explicit (cheap per step) but requires $h \lesssim 1/L$ for stability when $\nabla^2 V \preceq LI$, while the proximal scheme is unconditionally stable in h at the cost of solving an optimization problem each step.

Preconditioning via a non-Euclidean metric. The Euclidean norm in (10.3) is a choice. Replacing it by $\|y - x_k\|_M^2 := (y - x_k)^\top M (y - x_k)$ for some $M \succ 0$ produces the *preconditioned* proximal scheme

$$x_{k+1} = \arg \min_y \left[V(y) + \frac{1}{2h} \|y - x_k\|_M^2 \right], \quad (10.4)$$

with first-order condition $\nabla V(x_{k+1}) + \frac{1}{h} M(x_{k+1} - x_k) = 0$, i.e. $x_{k+1} = x_k - h M^{-1} \nabla V(x_{k+1})$. The continuous-time limit is the preconditioned flow

$$\dot{x}_t = -M^{-1} \nabla V(x_t). \quad (10.5)$$

Choosing $M = \nabla^2 V(x_k)$ locally recovers Newton's method; matching M to the global curvature of V removes anisotropy. The point worth recording is that the *metric in the proximal step determines the gradient flow*: a different metric gives a different flow, with the same equilibrium x^* but different trajectories and convergence constants.

Transport to probability space

We now transport this picture to sampling. The state is no longer a point $x \in \mathbb{R}^d$ but a probability measure $\mu \in \mathcal{P}(\mathbb{R}^d)$; the role of the objective V is played by the KL divergence to the target,

$$\mathcal{F}(\mu) := D_{\text{KL}}(\mu \parallel \pi), \quad \arg \min_{\mu} \mathcal{F}(\mu) = \pi, \quad (10.6)$$

attained uniquely at π where $\mathcal{F}(\pi) = 0$. The minimum encodes all information about the target.

For the proximal scheme to make sense on $\mathcal{P}(\mathbb{R}^d)$ we need a metric. The natural choice is the Wasserstein-2 distance W_2 (Theorem 4.3). With this choice, (10.3) becomes the *JKO scheme* [12]:

$$\mu_{k+1} = \arg \min_{\mu} \left[D_{\text{KL}}(\mu \parallel \pi) + \frac{1}{2h} W_2^2(\mu, \mu_k) \right]. \quad (10.7)$$

This is the probability-space analogue of (10.3), with $D_{\text{KL}}(\cdot \parallel \pi)$ in place of V and $W_2^2/2h$ in place of $\|\cdot - x_k\|^2/2h$.

Solving JKO infinitesimally. The aim is to identify a continuous-time flow as the $h \rightarrow 0$ limit of (10.7). We approximate both terms in (10.7) to first order in h around μ_k .

The minimizer μ_{k+1} in (10.7) is close to μ_k for small h , so we parametrize it as the pushforward of μ_k along an infinitesimal displacement field hv for some velocity $v : \mathbb{R}^d \rightarrow \mathbb{R}^d$:

$$\mu_{k+1} = (I + hv)_\# \mu_k. \quad (10.8)$$

The pair (μ_k, v) then sweeps out a curve $\mu_t^{(\text{lin})} = (I + tv)_\# \mu_k$ for $t \in [0, h]$ joining μ_k to μ_{k+1} , satisfying the continuity equation $\partial_t \mu_t^{(\text{lin})} + \nabla \cdot (\mu_t^{(\text{lin})} v) = 0$ at $t = 0$.

Approximating W_2^2 . The cost of moving every particle of μ_k along the displacement hv is $\mathbb{E}_{\mu_k}[\|hv(X)\|^2] = h^2 \int \|v\|^2 d\mu_k$. Since W_2^2 is the minimum transport cost between μ_k and μ_{k+1} , and the displacement (10.8) provides one admissible transport,

$$W_2^2(\mu_{k+1}, \mu_k) \leq h^2 \int \|v(x)\|^2 d\mu_k(x); \quad (10.9)$$

infinitesimally this inequality is an equality (the displacement is optimal to leading order), and we will treat it as such. The right-hand side is a quadratic form in v , the squared $L^2(\mu_k)$ norm; intuitively the kinetic energy needed to displace μ_k to μ_{k+1} .

Approximating $D_{\text{KL}}(\mu \parallel \pi)$. Differentiating $\mu_t^{(\text{lin})}$ at $t = 0$:

$$D_{\text{KL}}(\mu_{k+1} \parallel \pi) \approx D_{\text{KL}}(\mu_k \parallel \pi) + h \int (\log(\mu_k/\pi) + 1) \partial_t \mu_t^{(\text{lin})} \big|_{t=0} dx,$$

using $\delta D_{\text{KL}}(\mu \parallel \pi) = \int (\log(\mu/\pi) + 1) \delta \mu dx$. The perturbation is $\partial_t \mu_t^{(\text{lin})} \big|_{t=0} = -\nabla \cdot (\mu_k v)$. Integration by parts (the constant $+1$ drops since $\int \nabla \cdot (\mu_k v) dx = 0$):

$$D_{\text{KL}}(\mu_{k+1} \parallel \pi) \approx D_{\text{KL}}(\mu_k \parallel \pi) + h \int \nabla \log(\mu_k/\pi) \cdot v d\mu_k. \quad (10.10)$$

This is linear in v and identifies $\nabla \log(\mu_k/\pi)$ as the “gradient” of $D_{\text{KL}}(\cdot \parallel \pi)$ at μ_k in the $L^2(\mu_k)$ -pairing.

The quadratic problem. Substituting (10.9) and (10.10) into the JKO objective:

$$D_{\text{KL}}(\mu_{k+1} \parallel \pi) + \frac{1}{2h} W_2^2(\mu_{k+1}, \mu_k) \approx D_{\text{KL}}(\mu_k \parallel \pi) + h \int \nabla \log(\mu_k/\pi) \cdot v d\mu_k + \frac{h}{2} \int \|v\|^2 d\mu_k,$$

a strictly convex quadratic in v over $L^2(\mu_k)$. Pointwise stationarity in $v(x)$ gives

$$v(x) = -\nabla \log(\mu_k/\pi)(x). \quad (10.11)$$

From the proximal step to Fokker–Planck. Substituting (10.11) into the continuity equation and letting $h \rightarrow 0$ ($\mu_k \rightarrow \mu_t$, $\mu_{k+1} \rightarrow \mu_{t+dt}$):

$$\partial_t \mu_t + \nabla \cdot (\mu_t v_t) = 0, \quad v_t = -\nabla \log(\mu_t/\pi).$$

Writing $\pi \propto e^{-V}$ so $\nabla \log \pi = -\nabla V$, expanding $\nabla \log \mu_t = (\nabla \mu_t)/\mu_t$, and simplifying:

$$\partial_t \mu_t = \nabla \cdot (\mu_t \nabla V) + \Delta \mu_t, \quad (10.12)$$

which is the Fokker–Planck equation for the *Langevin SDE*

$$dX_t = -\nabla V(X_t) dt + \sqrt{2} dW_t. \quad (10.13)$$

Proposition 10.1 (Langevin = Wasserstein gradient flow). *For $\pi \propto e^{-V}$, the law μ_t of the Langevin SDE (10.13) is the small-step limit of the JKO scheme (10.7): the velocity field driving the continuity equation is the $L^2(\mu_t)$ -steepest-descent direction for $D_{\text{KL}}(\cdot \parallel \pi)$, namely $v_t = -\nabla \log(\mu_t/\pi)$.*

The construction parallels the Euclidean case term-by-term: proximal step (10.3) $\leftrightarrow$ JKO (10.7); quadratic penalty $\|y - x_k\|^2 \leftrightarrow$ infinitesimal transport cost $\int \|v\|^2 d\mu_k$ (10.9); linearization of V around $x_k \leftrightarrow$ linearization of $D_{\text{KL}}(\cdot \parallel \pi)$ around μ_k (10.10); minimizer $-\nabla V(x_k) \leftrightarrow$ minimizer $-\nabla \log(\mu_k/\pi)$ (10.11); flow $\dot{x} = -\nabla V \leftrightarrow$ Fokker–Planck (10.12).

Identifying the metric tensor. The parallel goes further: just as the Euclidean preconditioned flow (10.5) reads $\dot{x} = -M^{-1}\nabla V(x)$, the Wasserstein flow has the same structural form once we identify the metric tensor explicitly. Velocity fields and density perturbations are related by the continuity equation: a perturbation $\delta\mu$ generated by a potential-field velocity $v = \nabla\phi$ satisfies $\delta\mu = -\nabla \cdot (\mu \nabla \phi)$. Define the elliptic operator

$$\mathcal{L}_\mu := -\nabla \cdot (\mu \nabla \cdot), \quad (10.14)$$

acting on smooth scalar potentials. The map $\phi \mapsto \delta\mu = \mathcal{L}_\mu \phi$ is the bijection between potentials and tangent vectors at μ , and the W_2 inner product on velocities can be written as

$$\int \nabla \phi_1 \cdot \nabla \phi_2 d\mu = \int \phi_1 \mathcal{L}_\mu \phi_2 dx = \langle \phi_1, \mathcal{L}_\mu \phi_2 \rangle_{L^2},$$

showing that $\mathcal{L}_\mu$ is the W_2 metric tensor in the potential representation. The Fokker–Planck equation (10.12) can then be written as

$$\partial_t \mu_t = -\mathcal{L}_{\mu_t} [\log(\mu_t/\pi)], \quad (10.15)$$

where $\log(\mu_t/\pi)$ is the $L^2(dx)$ first variation of $D_{\text{KL}}(\cdot \parallel \pi)$ at μ_t . Compare with (10.5): $\dot{x} = -M^{-1}\nabla V$. The elliptic operator $\mathcal{L}_\mu$ plays the role of M^{-1} — the inverse metric — while $\log(\mu/\pi)$ plays the role of ∇V , the bare first variation. Both flows are “preconditioned” steepest descents on their objective, with the preconditioner determined by the metric: constant M on $\mathbb{R}^d$; the position-dependent $\mathcal{L}_\mu = -\nabla \cdot (\mu \nabla \cdot)$ on $\mathcal{P}(\mathbb{R}^d)$. The formal perspective is brought by Otto [17]. Discussions of other metric tensors one may refer to [6].

Convergence rate: log-Sobolev as Polyak–Łojasiewicz

In Euclidean optimization with α -strongly convex V , the Polyak–Łojasiewicz (PL) inequality

$$V(x) - V(x^*) \leq \frac{1}{2\alpha} \|\nabla V(x)\|^2 \quad (10.16)$$

holds and, combined with the dissipation identity $\frac{d}{dt} V(x_t) = -\|\nabla V(x_t)\|^2$ along (10.2), yields by Grönwall

$$V(x_t) - V(x^*) \leq e^{-2\alpha t} (V(x_0) - V(x^*)).$$

The sampling analogue is the *log-Sobolev inequality*: π satisfies LSI with constant $\alpha > 0$ if for all μ ,

$$D_{\text{KL}}(\mu \parallel \pi) \leq \frac{1}{2\alpha} \mathbb{E}_\mu[\|\nabla \log(\mu/\pi)\|^2]. \quad (10.17)$$

The right-hand side is $\frac{1}{2\alpha} \int \|v\|^2 d\mu$ with $v = \nabla \log(\mu/\pi)$ the JKO-velocity field (10.11), so LSI is exactly the PL inequality “ $\mathcal{F}(\mu) \leq \frac{1}{2\alpha} \|v\|_{L^2(\mu)}^2$ ” in the W_2 geometry. Combined with the dissipation identity along the WGF, $\frac{d}{dt} D_{\text{KL}}(\mu_t \parallel \pi) = -\mathbb{E}_{\mu_t}[\|\nabla \log(\mu_t/\pi)\|^2] = -\|v_t\|_{L^2(\mu_t)}^2$, LSI yields exponential decay:

$$D_{\text{KL}}(\mu_t \parallel \pi) \leq e^{-2\alpha t} D_{\text{KL}}(\mu_0 \parallel \pi). \quad (10.18)$$

When does π satisfy LSI? The Bakry–Émery criterion gives LSI with constant α whenever V is α -strongly convex — so e.g. Gaussians and strongly log-concave posteriors. Beyond strong convexity, LSI is preserved under bounded perturbations: if $\pi \propto e^{-V}$ satisfies LSI with constant α and $\pi' \propto e^{-(V+W)}$ with $\|W\|_\infty \leq C$, then π' satisfies LSI with constant αe^{-4C} . Thus mixtures and other mild non-log-concave perturbations still admit LSI, though with a constant that degrades with the perturbation size.

Worked example: Gaussian target. Take $\pi = \mathcal{N}(0, I)$ on $\mathbb{R}^d$, so $V(x) = \frac{1}{2}\|x\|^2$ is 1-strongly convex; this means $\alpha = 1$. Initialize at $\mu_0 = \mathcal{N}(m_0, \Sigma_0)$. The Langevin SDE preserves Gaussianity: $\mu_t = \mathcal{N}(m_t, \Sigma_t)$ with the moment ODEs

$$\dot{m}_t = -m_t, \quad \dot{\Sigma}_t = -2\Sigma_t + 2I, \quad (10.19)$$

obtained directly from Itô’s formula. Solving:

$$m_t = e^{-t} m_0, \quad \Sigma_t = e^{-2t} \Sigma_0 + (1 - e^{-2t}) I. \quad (10.20)$$

Mean decays at rate 1; covariance approaches I at rate 2. The KL between two Gaussians of equal mean is $D_{\text{KL}}(\mathcal{N}(m, \Sigma) \parallel \mathcal{N}(0, I)) = \frac{1}{2}(\text{tr} \Sigma - d - \log \det \Sigma + \|m\|^2)$; substituting (10.20) and expanding verifies $D_{\text{KL}}(\mu_t \parallel \pi) = O(e^{-2t})$, matching the LSI rate (10.18) with $\alpha = 1$.

For an anisotropic target $\pi = \mathcal{N}(0, \Lambda^{-1})$ with $\Lambda = \text{diag}(\lambda_1, \dots, \lambda_d)$, the dynamics decouple into d independent 1D systems. The i -th coordinate satisfies $\dot{m}_i = -\lambda_i m_i$, $\dot{\Sigma}_{ii} = -2\lambda_i \Sigma_{ii} + 2$, contracting at rate λ_i . The global mixing rate in W_2 is set by the slowest direction $\lambda_{\min}$: equilibration takes continuous time $\asymp 1/\lambda_{\min}$, and we will see in the next subsection that this becomes the condition number $\kappa = \lambda_{\max}/\lambda_{\min}$ in discrete iteration counts.

References

- [1] Michael Albergo, Nicholas M Boffi, and Eric Vanden-Eijnden. Stochastic interpolants: A unifying framework for flows and diffusions. *Journal of Machine Learning Research*, 26(209):1–80, 2025.

- [2] Michael Samuel Albergo and Eric Vanden-Eijnden. Building normalizing flows with stochastic interpolants. In *The Eleventh International Conference on Learning Representations*, 2023.
- [3] Brian DO Anderson. Reverse-time diffusion equation models. *Stochastic Processes and their Applications*, 12(3):313–326, 1982.
- [4] Sitan Chen, Sinho Chewi, Jerry Li, Yuanzhi Li, Adil Salim, and Anru R Zhang. Sampling is as easy as learning the score: theory for diffusion models with minimal data assumptions. In *International Conference on Learning Representations*, 2023.
- [5] Yifan Chen, Xiaoou Cheng, Jonathan Niles-Weed, and Jonathan Weare. Convergence of unadjusted langevin in high dimensions: Delocalization of bias. *Communications on Pure and Applied Mathematics*, page e70032, 2024.
- [6] Yifan Chen, Daniel Zhengyu Huang, Jiaoyang Huang, Sebastian Reich, and Andrew M Stuart. Sampling via gradient flows in the space of probability measures. *arXiv preprint arXiv:2310.03597*, 2023.
- [7] Sinho Chewi. Log-concave sampling. *Book draft available at <https://chewisinho.github.io>*, 9:17–18, 2023.
- [8] Hyungjin Chung, Jeongsol Kim, Michael Thompson Mccann, Marc Louis Klasky, and Jong Chul Ye. Diffusion posterior sampling for general noisy inverse problems. In *The Eleventh International Conference on Learning Representations*, 2023.
- [9] Prafulla Dhariwal and Alexander Nichol. Diffusion models beat gans on image synthesis. *Advances in neural information processing systems*, 34:8780–8794, 2021.
- [10] Carles Domingo-Enrich, Michal Drozdal, Brian Karrer, and Ricky TQ Chen. Ad-joint matching: Fine-tuning flow and diffusion generative models with memoryless stochastic optimal control. In *The Thirteenth International Conference on Learning Representations*.
- [11] Jonathan Ho and Tim Salimans. Classifier-free diffusion guidance. In *NeurIPS 2021 Workshop on Deep Generative Models and Downstream Applications*, 2021.
- [12] Richard Jordan, David Kinderlehrer, and Felix Otto. The variational formulation of the fokker–planck equation. *SIAM journal on mathematical analysis*, 29(1):1–17, 1998.
- [13] Ruilin Li, Hongyuan Zha, and Molei Tao. Sqrt (d) dimension dependence of langevin monte carlo. In *The International Conference on Learning Representations*, 2022.
- [14] Yaron Lipman, Ricky TQ Chen, Heli Ben-Hamu, Maximilian Nickel, and Matthew Le. Flow matching for generative modeling. In *The Eleventh International Conference on Learning Representations*, 2023.

- [15] Jun S Liu and Jun S Liu. *Monte Carlo strategies in scientific computing*, volume 10. Springer, 2001.
- [16] Xingchao Liu, Chengyue Gong, and Qiang Liu. Flow straight and fast: Learning to generate and transfer data with rectified flow. In *The Eleventh International Conference on Learning Representations (ICLR)*, 2023.
- [17] F Otto. The geometry of dissipative evolution equations: The porous medium equation. *Communications in Partial Differential Equations*, 26(1-2):101–174, 2001.
- [18] Daniel Sanz-Alonso and Omar Al-Ghattas. A first course in monte carlo methods. *arXiv preprint arXiv:2405.16359*, 2024.
- [19] Daniel Sanz-Alonso, Andrew Stuart, and Armeen Taeb. *Inverse problems and data assimilation*, volume 107. Cambridge University Press, 2023.
- [20] Yuyang Shi, Valentin De Bortoli, Andrew Campbell, and Arnaud Doucet. Diffusion schrödinger bridge matching. *Advances in neural information processing systems*, 36:62183–62223, 2023.
- [21] Yang Song and Stefano Ermon. Generative modeling by estimating gradients of the data distribution. *Advances in neural information processing systems*, 32, 2019.
- [22] Yang Song, Jascha Sohl-Dickstein, Diederik P Kingma, Abhishek Kumar, Stefano Ermon, and Ben Poole. Score-based generative modeling through stochastic differential equations. In *International Conference on Learning Representations*, 2021.
- [23] Wenpin Tang and Hanyang Zhao. Score-based diffusion models via stochastic differential equations. *Statistic Surveys*, 19:28–64, 2025.